

Databases

Database Design Using Entity-Relationship Model

João R. Campos

Bachelor in Informatics Engineering
Department of Informatics Engineering
University of Coimbra
2024/2025





From Previous Lesson(s)...

- Relational Model
 - Relational databases
 - Relation or table, tuple or row, attribute or column
 - Superkeys, candidate keys, primary key and foreign key
 - Integrity restrictions: primary key (entity), foreign key (referential), domain, ...
- SQL
 - create table, alter table, and drop table
 - Query the database: select...
 - Cartesian product and joining tables
 - Aggregation functions
 - Modifying the data (insert, update, delete)



Outline

- Database Design Process
- Entity-Relationship Model
 - Entities and Entity Sets
 - Relationships and Relationship Sets
 - Mapping Cardinalities and Participation
 - Removing Redundant Attributes
- From E-R Diagrams to Relational Schemas
 - Obtaining Tables from Entity Sets and Relationship Sets
 - The *onda* Tool
 - Obtaining the SQL DDL commands

Register your presence at UCStudent!

*These slides use the following book as reference:
Abraham Silberschatz, Henry F. Korth and S. Sudarshan,
“Database System Concepts”, McGraw-Hill Education,
Seventh Edition, 2019.*

This class focuses mostly on **Chapter 6**



Database Design

- Creating a database application is a complex task, involving several aspects, such as:
 - Design of the database schema
 - Implementation of the programs that access and update the data
 - Design of a security scheme to control access to data
- For small applications, we may be able to decide directly on the relations that are needed (e.g., the employee database seen before)
- For complex applications, we need to follow a process in order to obtain the best model relational schema possible



Database Design Process

- Characterize the **data needs** of the prospective database users
- **Conceptual design**
 - Choose a modeling approach (e.g. E-R model, normalization) and apply the concepts of the chosen data model
 - Translate requirements into a conceptual schema (e.g. E-R diagram)
- Specification of **functional requirements**
 - Conceptual schema can be used to extract the functional requirements
 - Describe the kinds of operations (or transactions) that will be performed on the data
- Implementation of the database
 - **Logical design**: from the conceptual schema to the relational data model
 - **Physical design**: decision related to the the data model (e.g., indexes needed)



Design Alternatives

- There may be several design alternatives for the same problem
- We must ensure that we avoid two major pitfalls:
 - **Redundancy**: a bad design may result in repeated information
 - **Incompleteness**: a bad design may make certain aspects of the enterprise difficult or impossible to model
- Avoiding bad designs is not enough!
 - There may be a number of good designs from which we must choose

why is this
a problem?



Design Approaches

- Entity-Relationship model
 - Models an organization as a collection of entities and relationships
 - **Entity**: a “thing” or “object” in the enterprise that is distinguishable from other objects, and is described by a set of attributes
 - **Relationship**: an association among several entities
 - Represented diagrammatically by an **entity-relationship diagram**
- Normalization theory
 - Formalize what designs are bad, and test for them

Databases

Entities and Entity Sets

Relationships and Relationship Sets

Mapping Cardinalities and Participation

Removing Redundant Attributes

ENTITY-RELATIONSHIP MODEL



E-R Data Model

- Developed to facilitate database design by allowing specification of a schema that represents the overall logical structure of a database
- The E-R model is very useful in mapping real-world meanings and interactions onto a **conceptual schema**
- Employs three basic concepts:
 - Entity sets
 - Relationship sets
 - Attributes
- Has an associated diagrammatic representation, the **E-R diagram**
- Many tools based on the E-R model are available



Entities and Entity Sets

- An **entity** is a “thing” or “object” in the real world that is distinguishable from all other objects
 - e.g., each person in a university is an entity
- An entity has a set of **properties**, and the values for some set of properties must uniquely identify an entity
 - e.g., a person may have a person *id* property whose value uniquely identifies that person
 - The value 2018280021 for person *id* would uniquely identify one particular person in the university
- An **entity set** is a set of entities of the same type that share the same properties or attributes
 - The set of all people who are instructors can be defined as the entity set instructor



Attributes and Values

- Entity sets do not need to be disjoint
 - e.g., it is possible to define the entity set person consisting of all people in a university
 - A person entity may be an instructor entity, a student entity, both, or neither
- An entity is represented by a set of **attributes**
 - Descriptive properties possessed by each member of an entity set
 - e.g., possible attributes of the instructor entity set are *ID*, *name*, *dept_name*, and *salary*
 - Attributes may be simple or complex (composed by sub-attributes)
- Each entity has a **value** for each of its attributes
 - e.g., a particular instructor entity may have the value 12121 for ID, the value Wu for name, the value Finance for dept name, and the value 90000 for salary

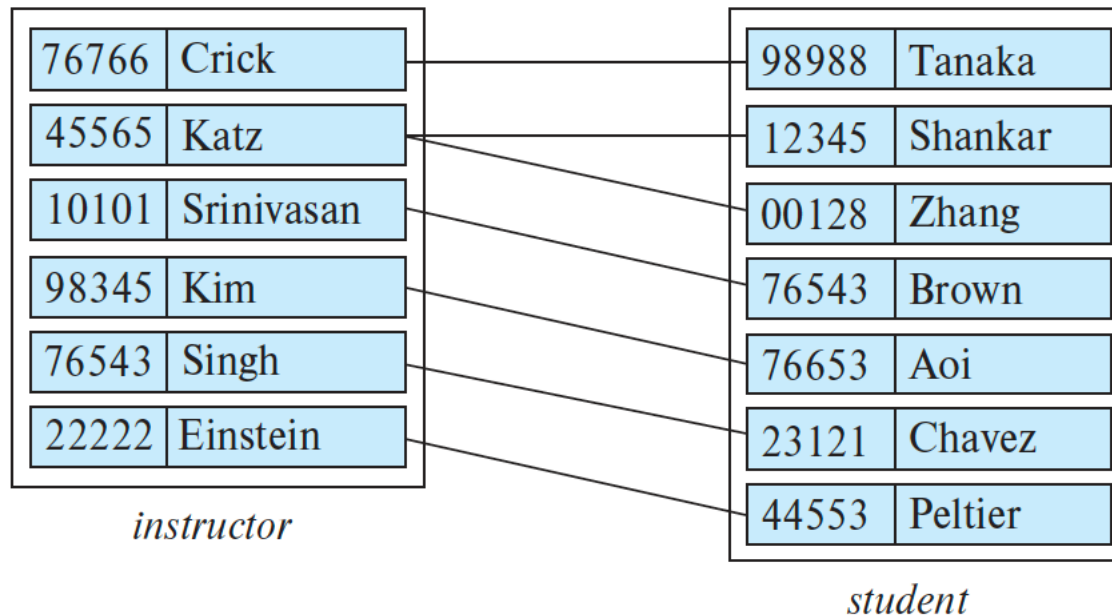
Entity Sets in the E-R Diagram...

instructor		
ID	VChr	PK
name	VChr	
salary	Number	

student		
ID	VChr	PK
name	VChr	
tot_cred	Number	

Relationships and Relationship Sets

- A **relationship** is an association among several entities
 - e.g., we can define a relationship *advisor* that associates instructor Katz with student Shankar, specifying that Katz is an advisor to student Shankar
- A **relationship set** is a set of relationships of the same type
 - Consider entity sets *instructor* and *student*: the relationship set *advisor* denotes the associations between students and the instructors who act as their advisors



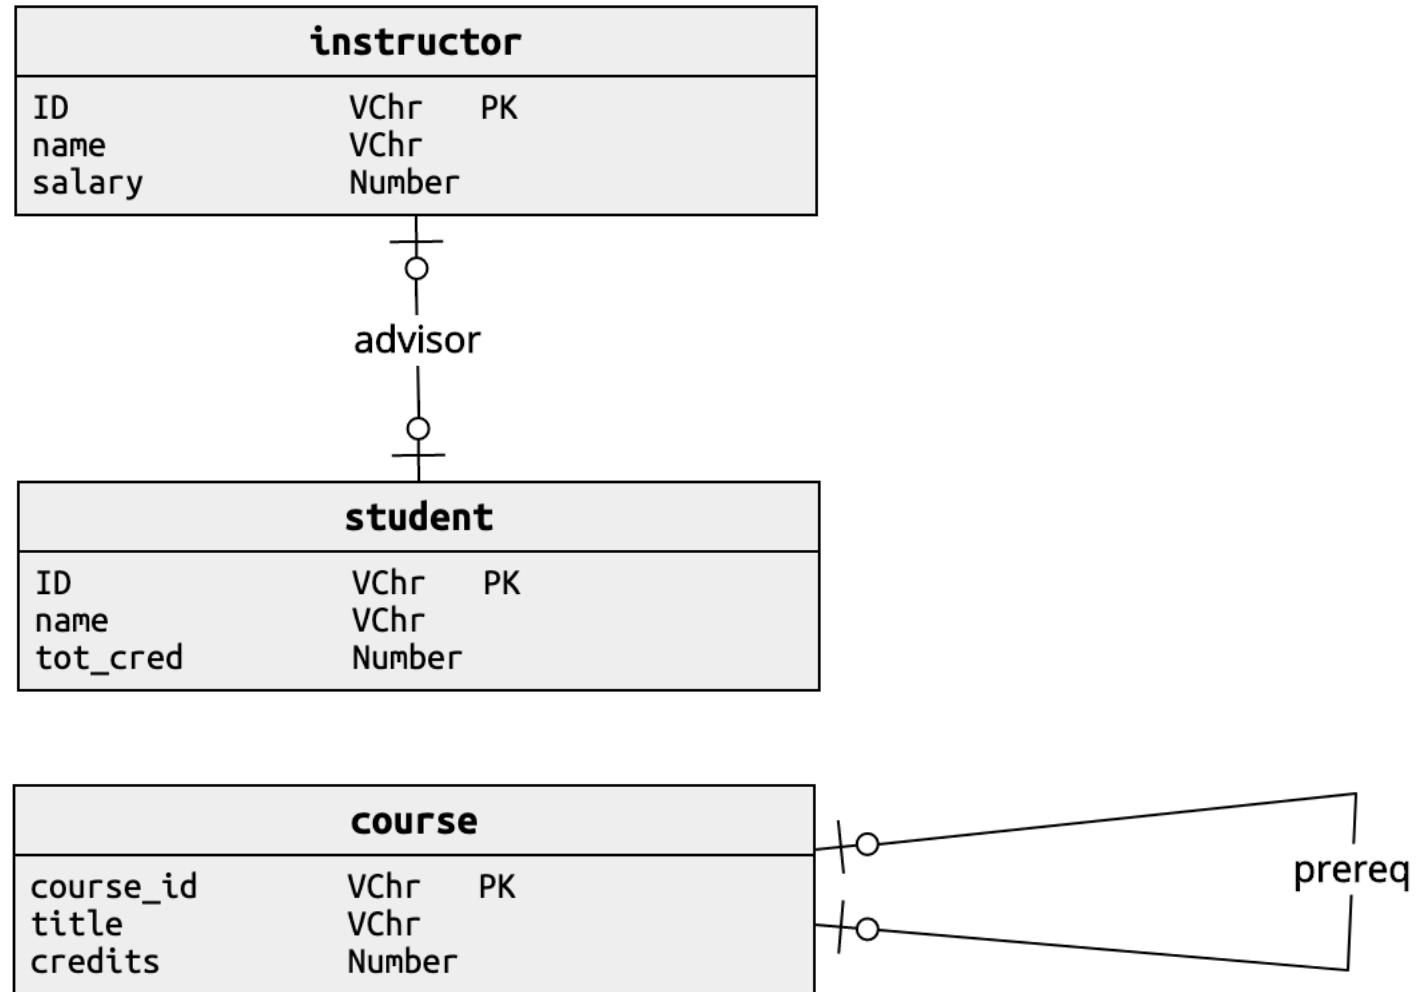
Source: A. Silberschatz, H. F. Korth and S. Sudarshan, "Database System Concepts", McGraw-Hill Education, Seventh Edition, 2019.



Relationships and Relationship Sets

- A **relationship instance** represents an association between the named entities in the real-world organization that is being modeled
 - e.g., the individual *instructor* entity Katz and the *student* entity Shankar participate in a relationship instance of *advisor*
- There may be several relationship sets between two entity sets
 - e.g., besides advising a student in the context of an internship, an instructor may also act as a tutor in a more general context
 - Relationship set *tutor* is thus different from relationship set *advisor*
- Relationship sets may be **recursive** when the same entity set participates more than once in the same relationship set

Relationship Sets in the E-R Diagram...




Attributes

- For each attribute, there is a set of permitted values, called the **domain, or value set**, of that attribute
 - The domain of attribute *course_id* might be the set of all text strings of a certain length
 - The domain of attribute *semester* might be strings from the set: *{Fall, Winter, Spring, Summer}*
- An attribute can be characterized by the following **attribute types**:
 - Simple attributes (e.g., *student_ID*) and composite attributes (e.g., *address*)
 - Single-valued (e.g., *student_ID*) and multivalued (e.g., *dependent_names*)
 - Derived attributes
 - e.g. age can be derived from the date of birth; should it be stored?
- An attribute takes a **null** value when an entity does not have a value for it


Mapping Cardinalities

- **Mapping cardinalities** express the number of entities to which another entity can be associated via a relationship set
- For a binary relationship set R between entity sets A and B , the mapping cardinality must be one of the following:

- One-to-one
- One-to-many
- Many-to-one
- Many-to-many




one instructor
one student




one instructor
many students

Source: A. Silberschatz, H. F. Korth and S. Sudarshan, "Database System Concepts", McGraw-Hill Education, Seventh Edition, 2019.

Mapping Cardinalities



many instructors
one student

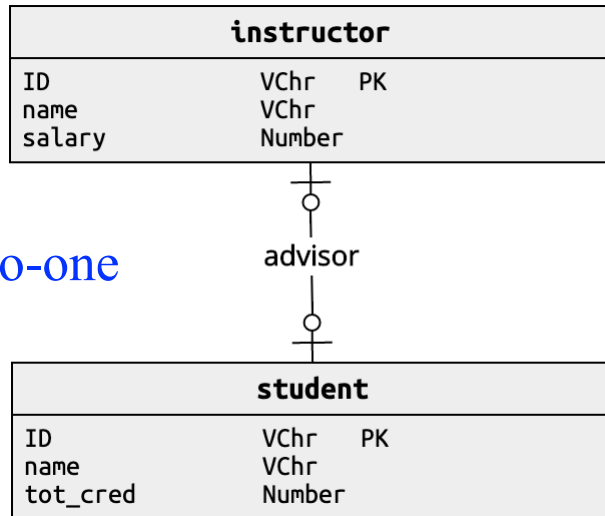


many instructors
many students

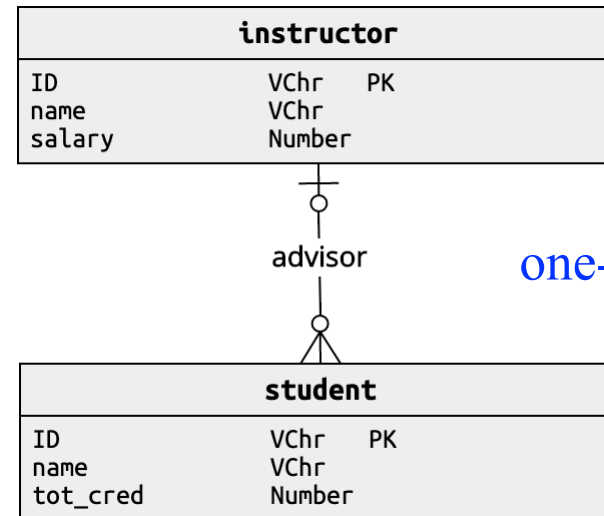
Source: A. Silberschatz, H. F. Korth and S. Sudarshan, "Database System Concepts", McGraw-Hill Education, Seventh Edition, 2019.

Cardinalities in the E-R Diagram...

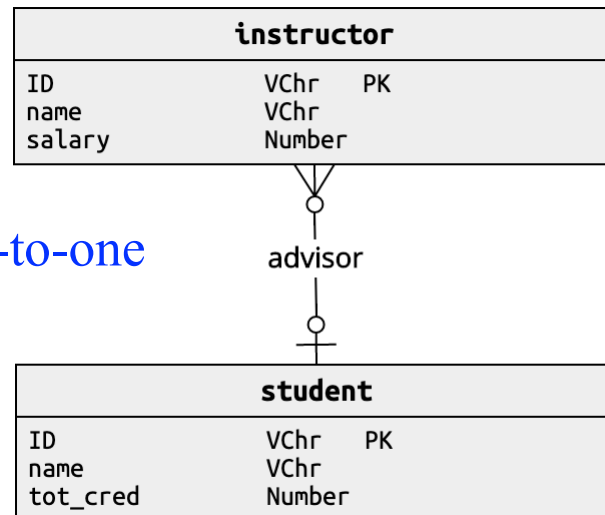
one-to-one



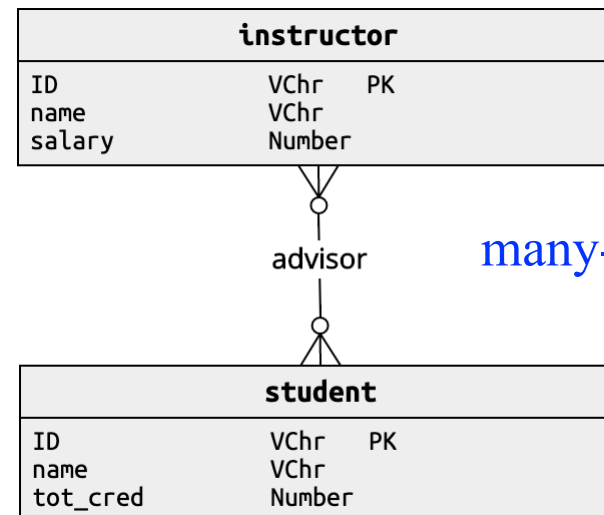
one-to-many



many-to-one

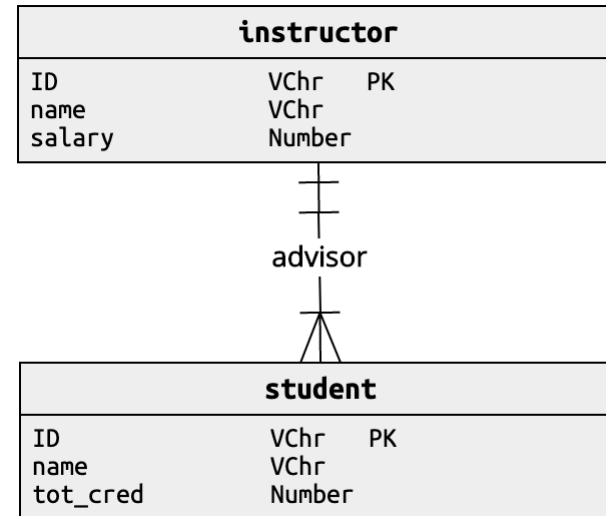
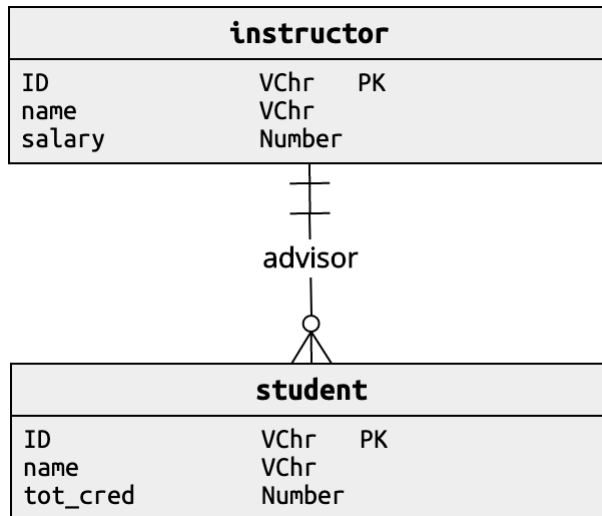


many-to-many



Total and Partial Participation

- **Total participation**: every entity in the entity set E must participate in at least one relationship in relationship set R
- **Partial participation**: if it is possible that some entities in E do not participate in relationships in R



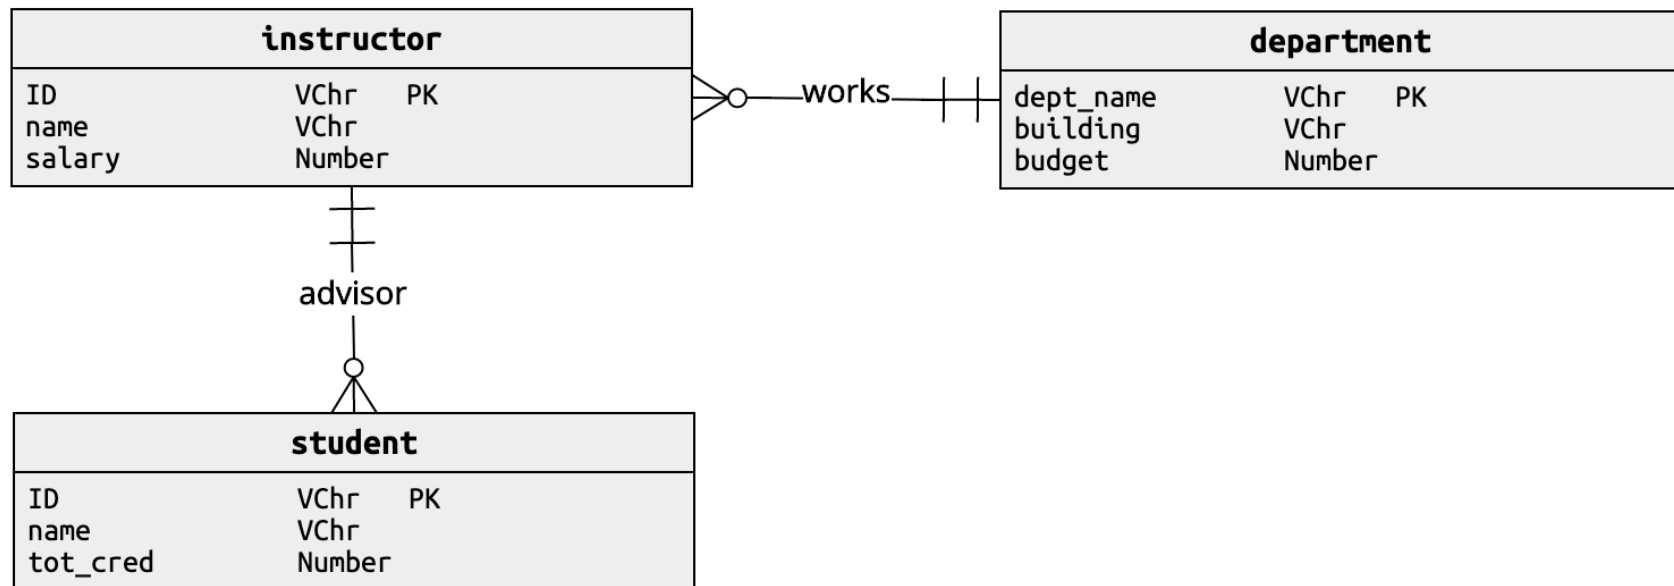


Defining the Primary Key

- The **primary key** is a set of attributes whose values univocally allow identifying a specific entity in the context of the organization
 - e.g., no two students have the same *student id* at the university
- Primary key and candidate key concepts from relations apply here in the same way
- How to select the primary key?
- Identical to the principles of relational PK discussed in previous class

Multiple Binary Relationships

- The same entity set may have several relationships sets with several other entity sets
 - e.g., entity set A may have a relationship set R_1 with entity set B and a relationship R_2 with entity set C



Removing Redundant Attributes

- A good E-R design does not include redundant attributes

instructor		
ID	VChr	PK
name	VChr	
dept_name	VChr	
salary	Number	

Removing Redundant Attributes



- A good E-R design does not include redundant attributes

instructor		
ID	VChr	PK
name	VChr	
dept_name	VChr	
salary	Number	

department		
dept_name	VChr	PK
building	VChr	
budget	Number	

Removing Redundant Attributes


- A good E-R design does not include redundant attributes





Alternative Notations for E-R Diagrams

- There exist several notations for E-R Diagrams
- We are using the **Crow's Foot Notation**
- All notations are based in the concepts presented before



entity set E with simple attribute A1, composite attribute A2, multivalued attribute A3, derived attribute A4, and primary key A1





many-to-many relationship





one-to-one relationship



many-to-one relationship




participation in R: total (E1) and partial (E2)



weak entity set



generalization



total generalization



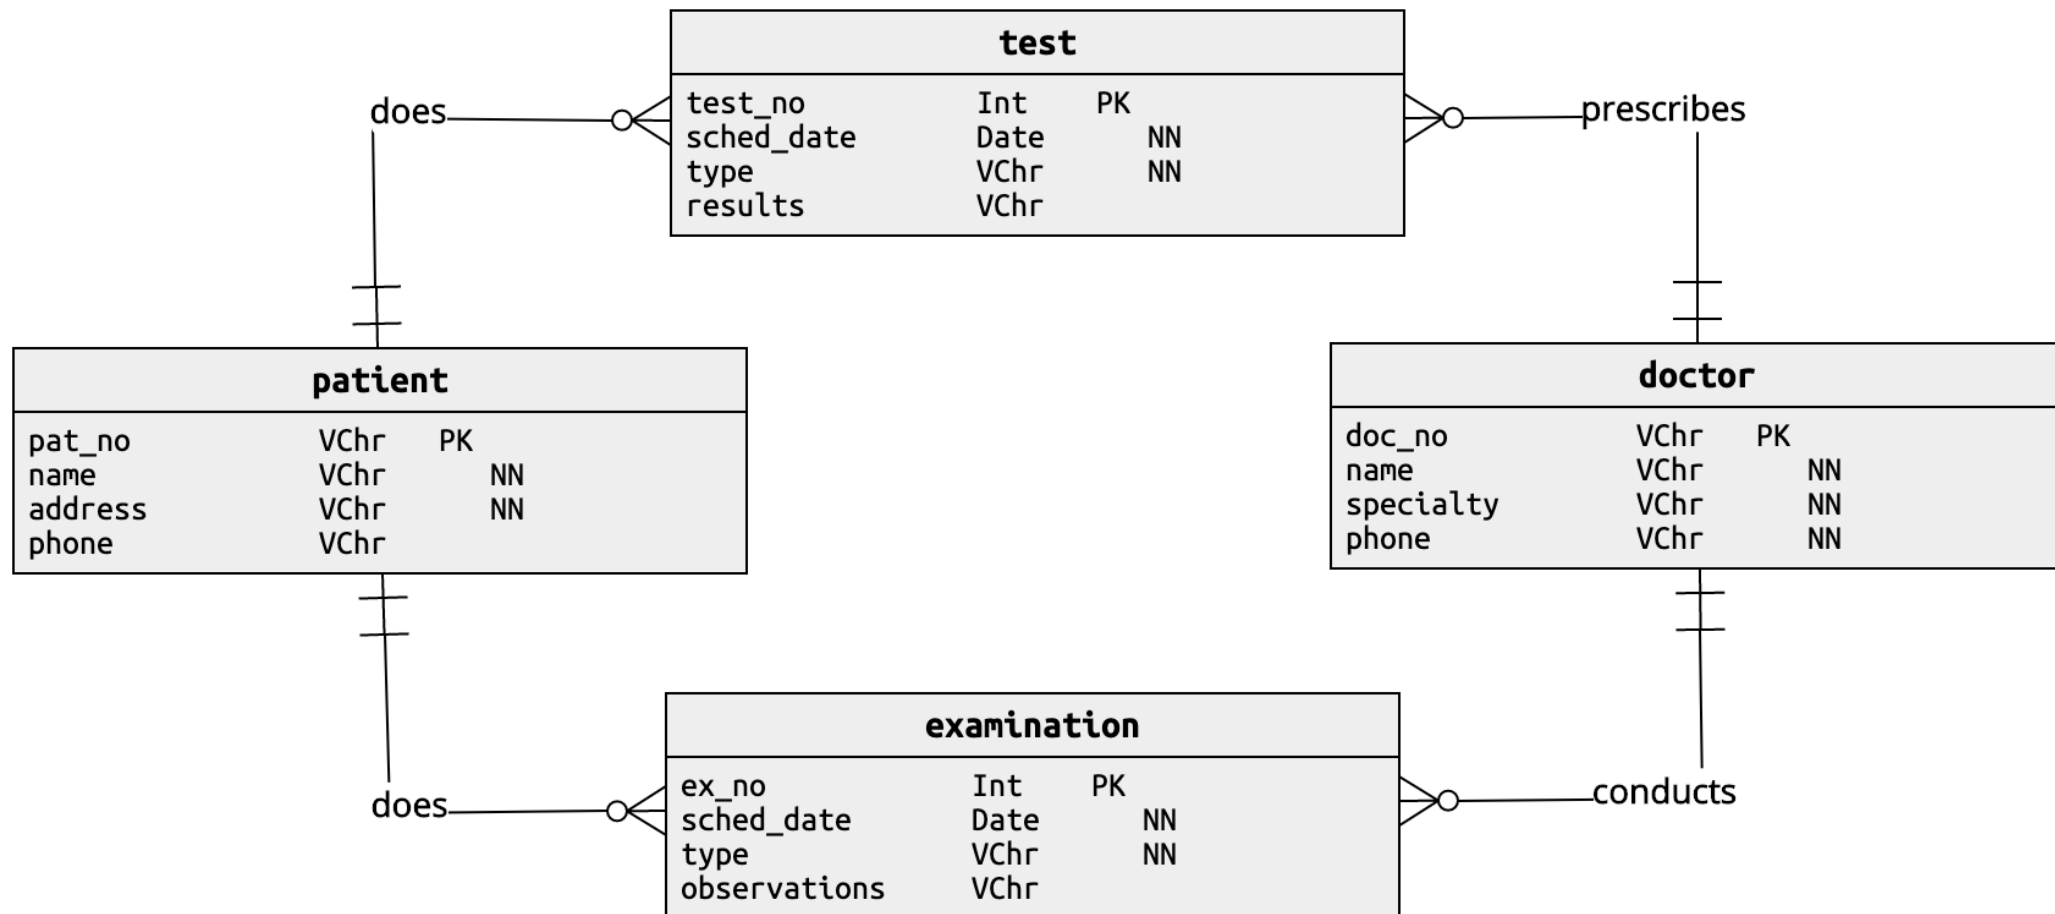
Source: A. Silberschatz, H. F. Korth and S. Sudarshan, "Database System Concepts", McGraw-Hill Education, Seventh Edition, 2019.

DEMO #1



- Assume the following:
 - *A small hospital needs to develop a new database application to manage patients. Over time, the hospital treats many patients and has a group of medical doctors. For each patient, the hospital wants a log of the various tests and examinations conducted. Tests are prescribed by doctors, while examinations are conducted by doctors.*
- TODO:
 - Entity sets?
 - Attributes of each entity?
 - Primary keys?
 - Relationship sets?
 - Cardinalities and participations?
 - E-R diagram...

DEMO #1 – Potential E-R Diagram

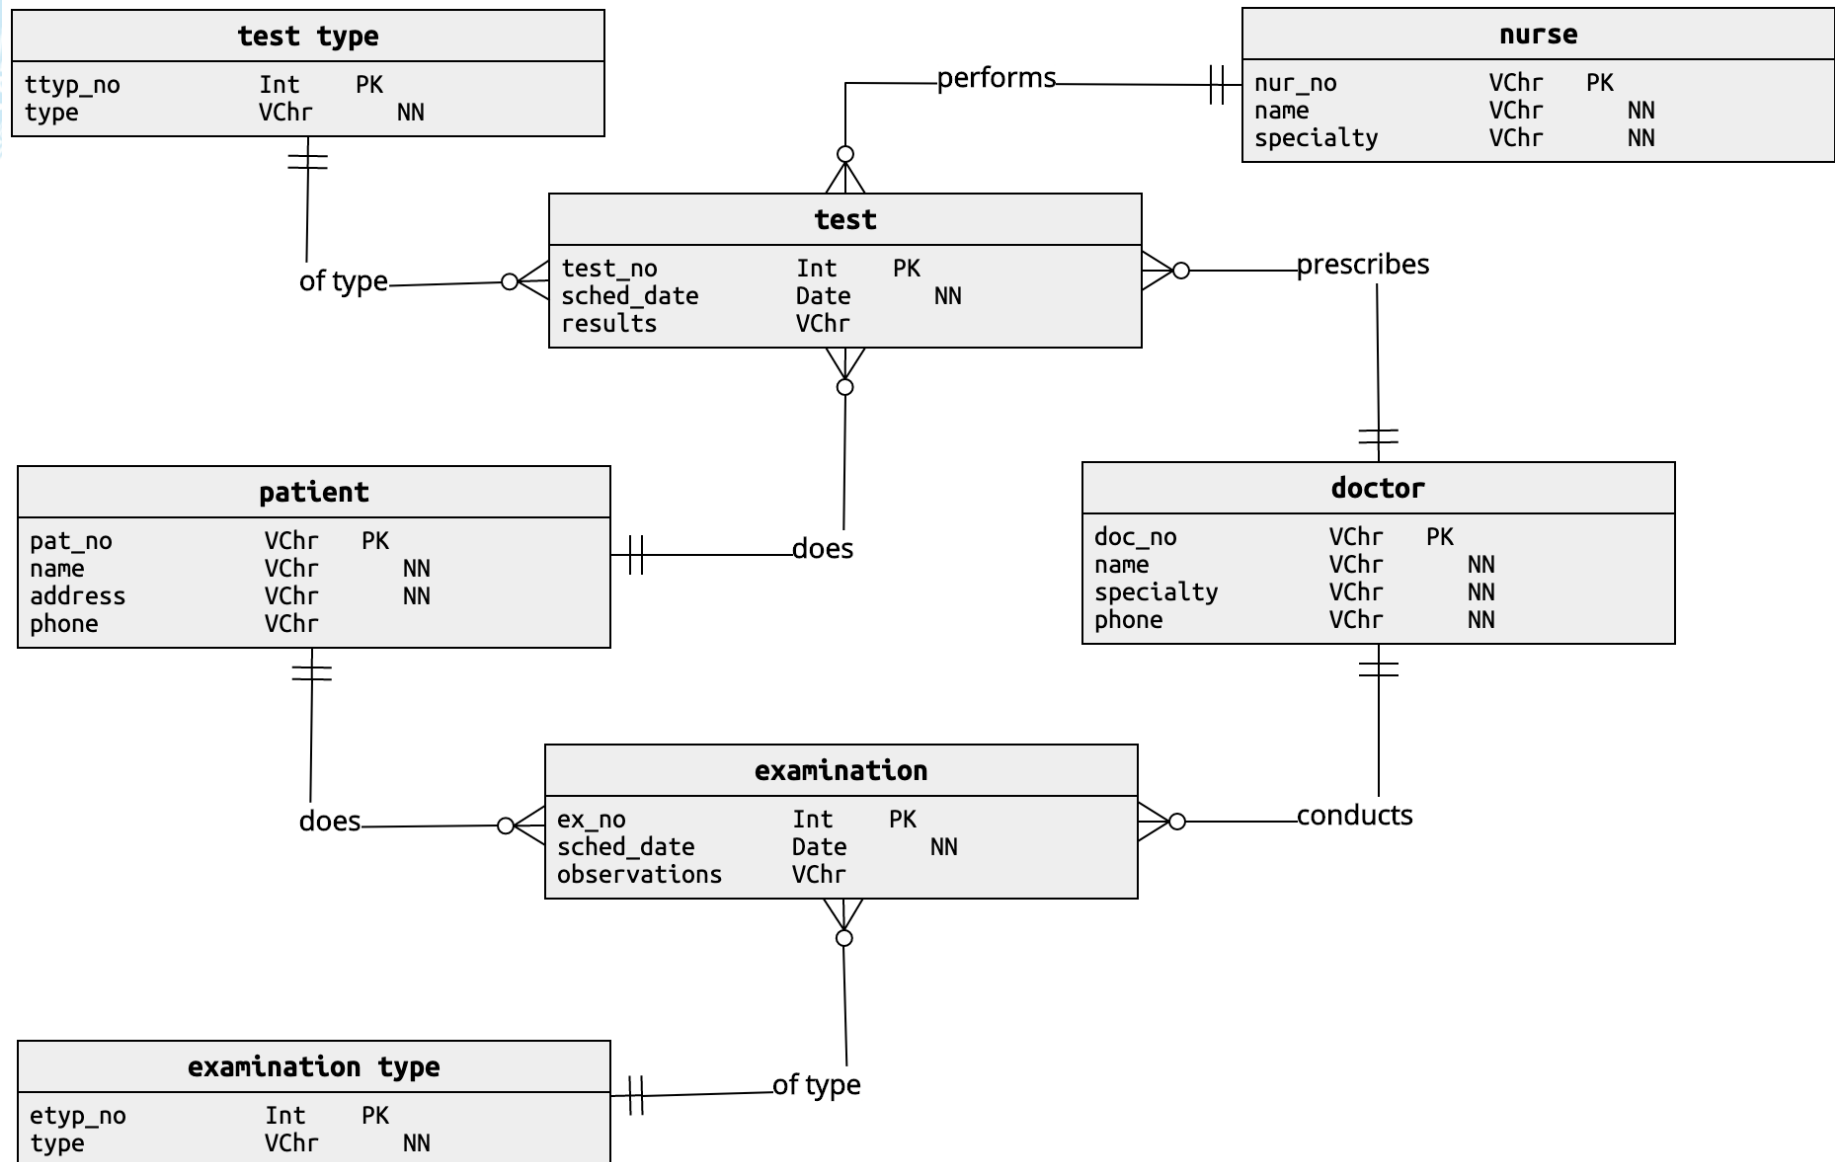


DEMO #2



- Assume the following:
 - *A small hospital needs to develop a new database application to manage patients. Over time, the hospital treats many patients and has a group of medical doctors. For each patient, the hospital wants a log of the various tests and examinations conducted. Tests are prescribed by doctors, while examinations are conducted by doctors.*
- TODO:
 - Improve the E-R diagram from DEMO #1 considering:
 - Tests are performed by nurses
 - There is a predefined set of types of tests, which can evolve over time
 - There is a predefined set of types of examinations, which can evolve over time
 - Implement the diagram using the *onda* tool
 - <http://onda.dei.uc.pt> (v4)

DEMO #2 – Potential E-R Diagram



Databases

Obtaining Tables from Entity Sets and Relationship Sets

The ONDA Tool

Obtaining the SQL DDL commands

**FROM E-R DIAGRAMS TO
RELATIONAL SCHEMAS**

The Goal

- Starting from the E-R Diagram, obtain the logical database design – relational schema
 - Relations, attributes, keys, constraints, etc.
 - Foreign-keys are generated during this process!
 - As seen before, no foreign-keys are represented in E-R diagrams (only relationship sets exist)
- Logical design may need to be fine-tuned (e.g., using normalization)
- The physical database design is built on top of this logical design
- SQL DDL commands can be automatically generated from the logical design in order to create the database
 - The SQL commands generated depend on the target database engine, as different engines implement slightly different versions of SQL

Strong Entity Sets \rightarrow Relations

- Let E be a strong entity set with attributes a_1, a_2, \dots, a_n
- This entity generates a relation schema E with n distinct attributes
 - Each tuple in relation E corresponds to one entity of the entity set E
- The primary key of the entity set serves as the primary key of the relation schema
 - This follows directly from the fact that each tuple corresponds to a specific entity in the entity set

entity set

department		
dept_name	VChr	PK
building	VChr	
budget	Number	

relation schema

department		
dept_name	VChr	PK
building	VChr	
budget	Numeric	



Relationship Sets \rightarrow Relations


- Each strong entity set in the relationship set will lead to a relation
 - (there is one exception, which we will discuss next class)
- Depending on the mapping cardinalities and participation, several cases may occur
- An **additional relation schema** may (or may not) be needed
 - Primary key and foreign keys must be defined
- A **foreign key** may (or may not) be added to one of the relations originated by the strong entity sets
- Both foreign keys and new relations are needed to allow setting the relationships between tuples
 - We should avoid situations where foreign keys may have *null* values!

One-to-One Relationship Sets: Case #1

- Let's assume a **one-to-one relationship set** with **partial participation** on one entity set and **full participation** on the other




- What is the solution? We have the same three hypothesis:
 - Relation schema A includes a foreign key to relation schema B
 - Relation schema B includes a foreign key to relation schema A
 - A relation A_B is originated to allow the mapping between A and B



Why is not the third hypothesis a good one?

One-to-One Relationship Sets: Case #2

- Let's now assume a **one-to-one relationship set** with **partial participation on both entity sets**



- Each entity set leads to a relation schema


a		
id	Int	PK

b		
id	Int	PK

- What about the relationship set? What does it originate?

One-to-One Relationship Sets: Case #2


- **First hypothesis:** relation schema A includes a foreign key to relation schema B



- Is this a good solution? Will there be tuples in relation A where the foreign key b_id is *null*?
 - Yes! According to the E-R diagram, there is partial participation, so some A entities do not have a correspondence to a B entity

One-to-One Relationship Sets: Case #2


- **Second hypothesis:** relation schema B includes a foreign key to relation schema A



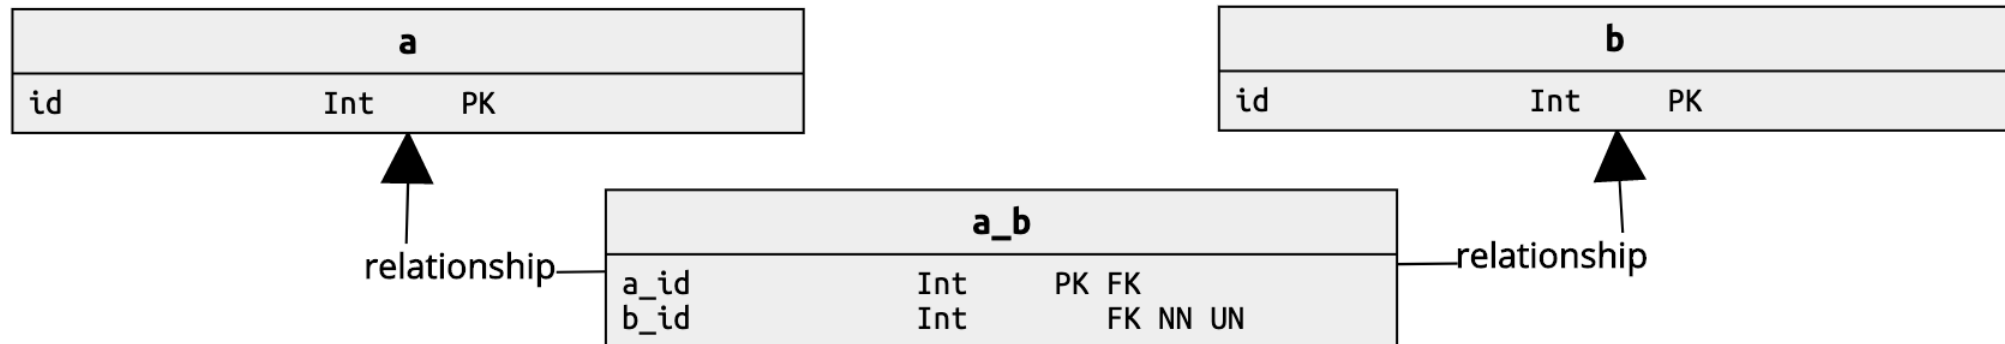
- Is this a good solution? Will there be tuples in relation B where the foreign key a_id is *null*?
 - Yes! According to the E-R diagram, there is partial participation, so some B entities do not have a correspondence to a A entity

One-to-One Relationship Sets: Case #2

- **Third hypothesis:** a relation A_B is originated to allow the mapping



- What are the attributes of the new relation schema A_B ?
- What is the primary key of A_B ?
- Are there any foreign keys?



One-to-One Relationship Sets: Case #3

- Finally, let's assume a **one-to-one relationship set** with **full participation on both entity sets**

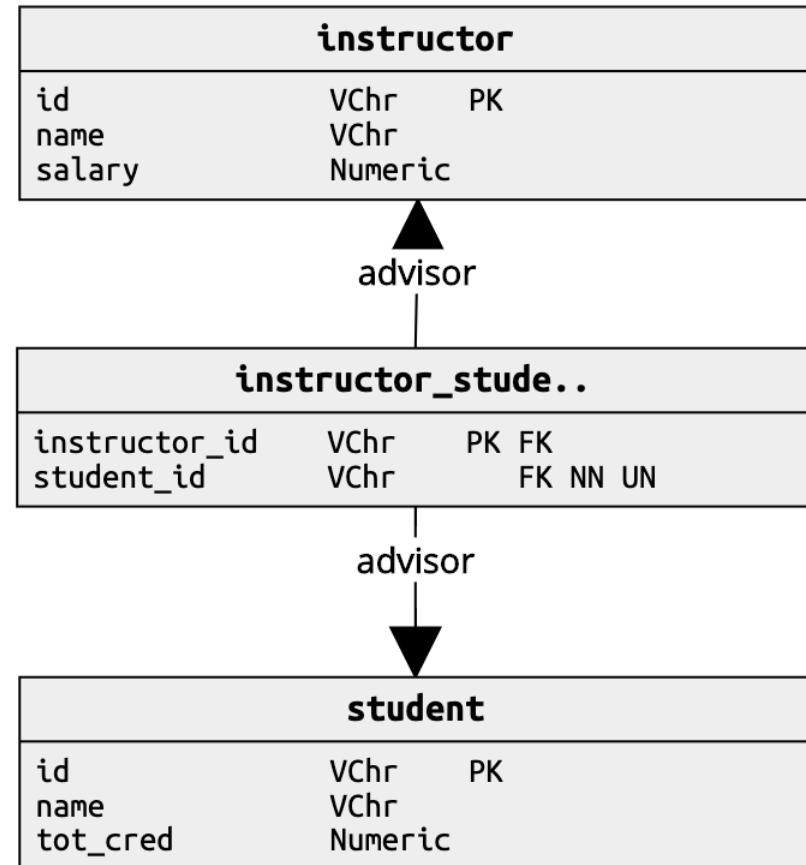
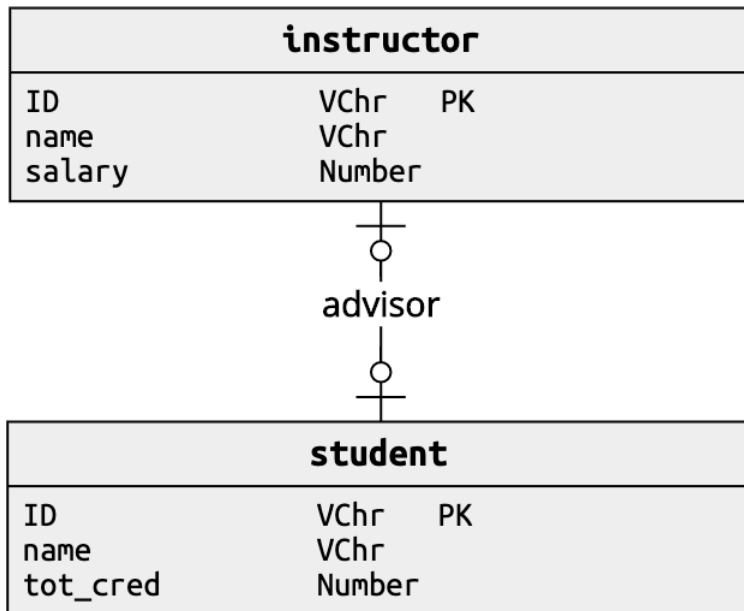


- What is the solution? We have the same three hypothesis:
 - Relation schema A includes a foreign key to relation schema B
 - Relation schema B includes a foreign key to relation schema A
 - A relation A_B is originated to allow the mapping between A and B

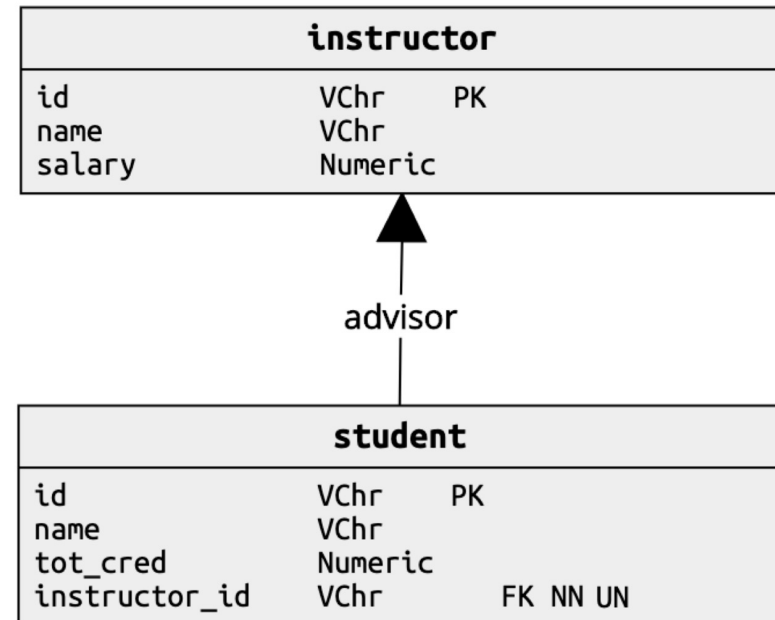
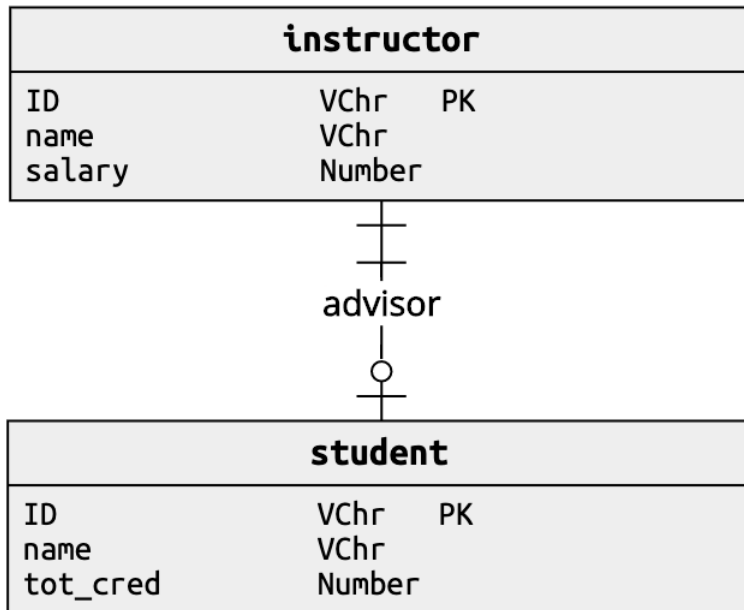
None of these
is good!

a_b				
id	Int	PK		
b_id	Int		NN	UN

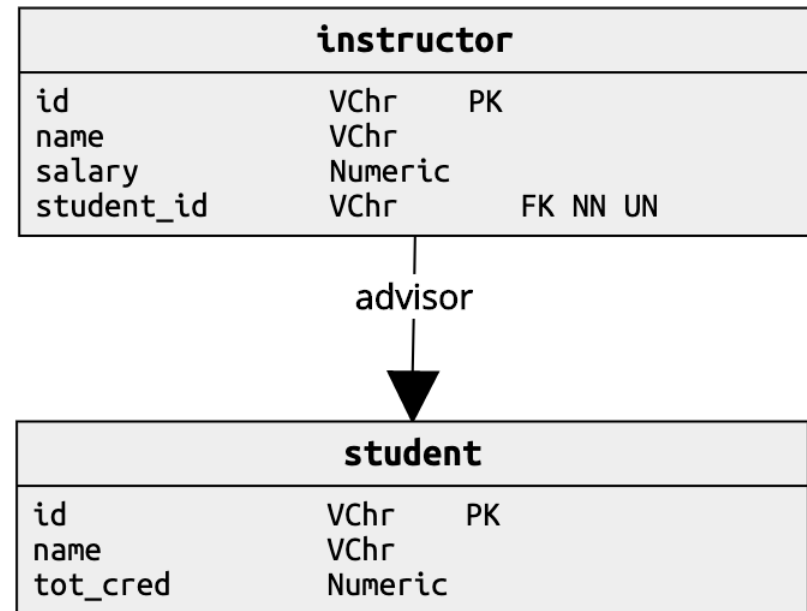
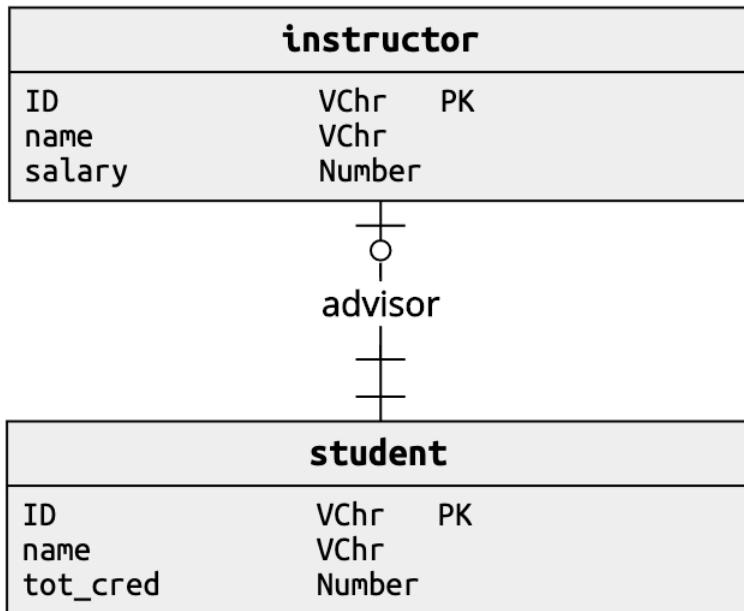
One-to-One Relationship Sets: Examples



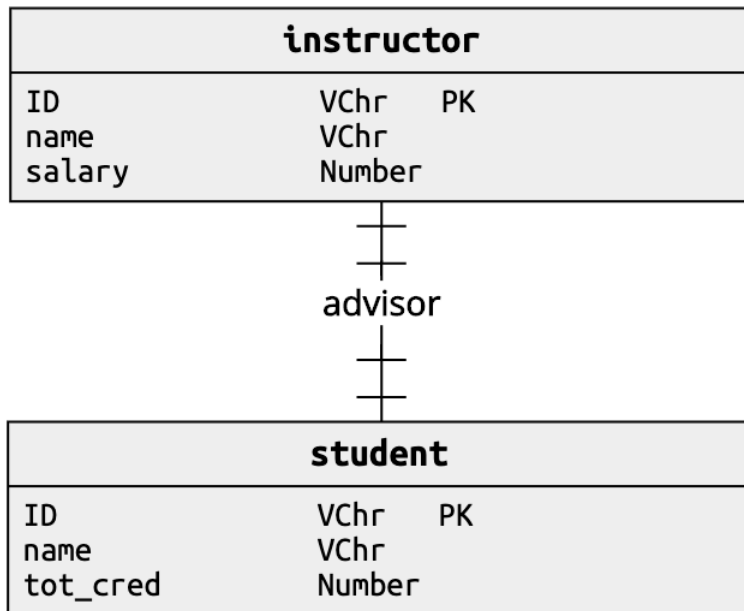
One-to-One Relationship Sets: Examples



One-to-One Relationship Sets: Examples




One-to-One Relationship Sets: Examples




instructor_stude..			
id	VChr	PK	
name	VChr		
salary	Numeric		
student_id	VChr		NN UN
student_name	VChr		
student_tot_cred	Numeric		

One-to-Many Relationship Sets: Case #1

- Let's assume a **one-to-many relationship set** with **total participation on the *many* side**




- What is the solution?




One-to-Many Relationship Sets: Case #2

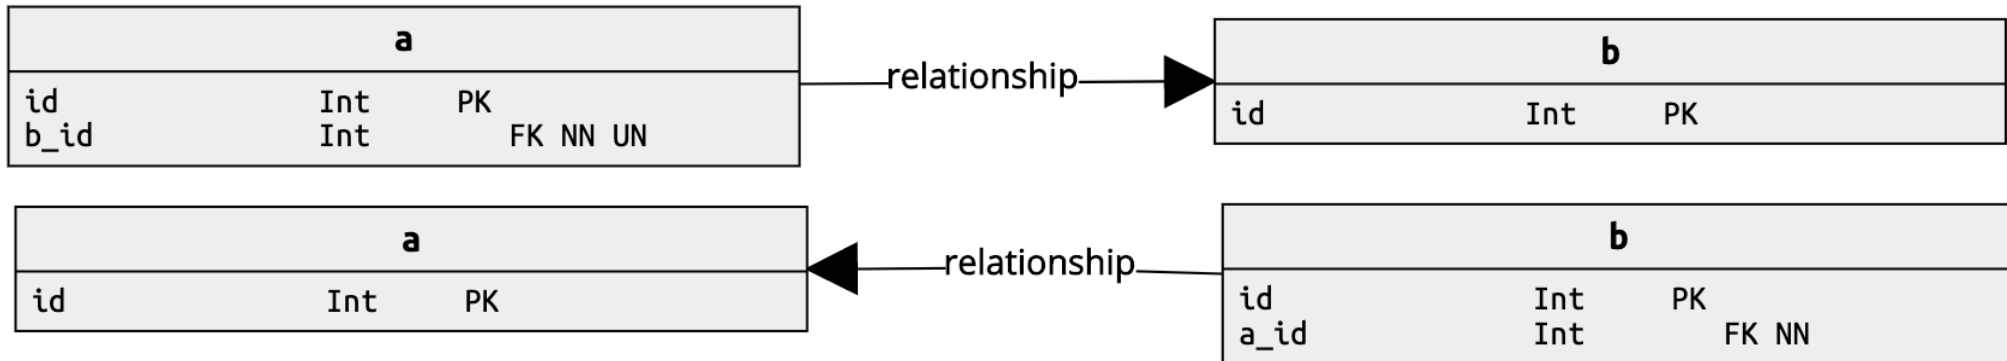
- Let's assume a **one-to-many relationship set** with **partial participation** on both entity sets



- Can this be represented with a single relation schema?




- What about with two tables?

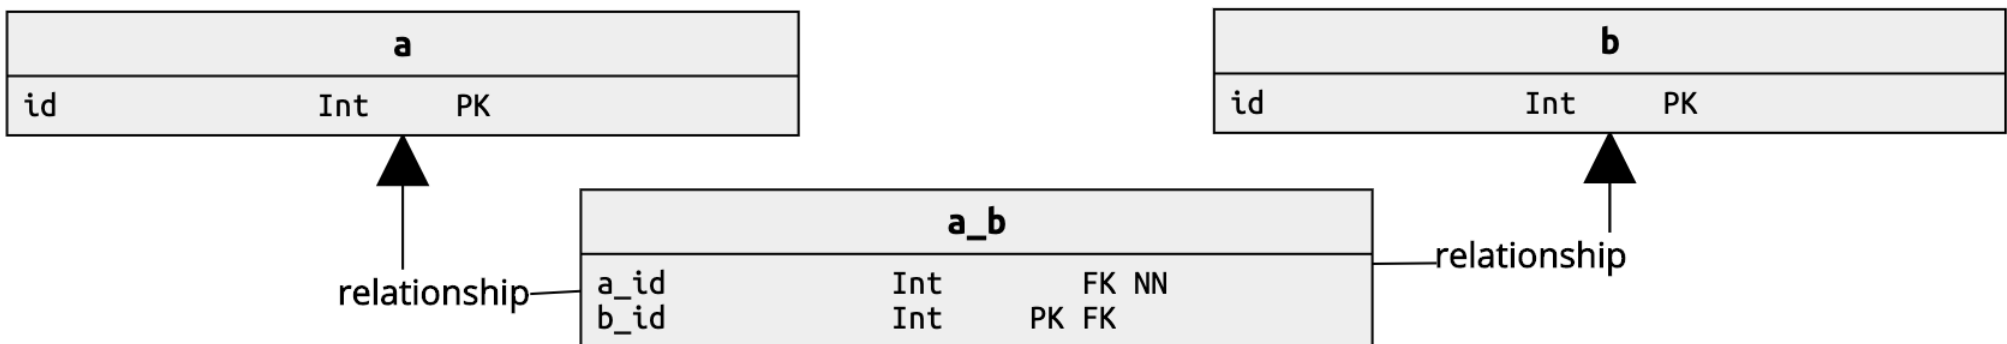


One-to-Many Relationship Sets: Case #2

- What about three tables?



- What are the attributes of the new relation schema *A_B*?
- What is the primary key of *A_B*?
- Are there any foreign keys?




One-to-Many Relationship Sets: Case #3

- Finally, let's assume a **one-to-many relationship set** with **total participation on both sides**

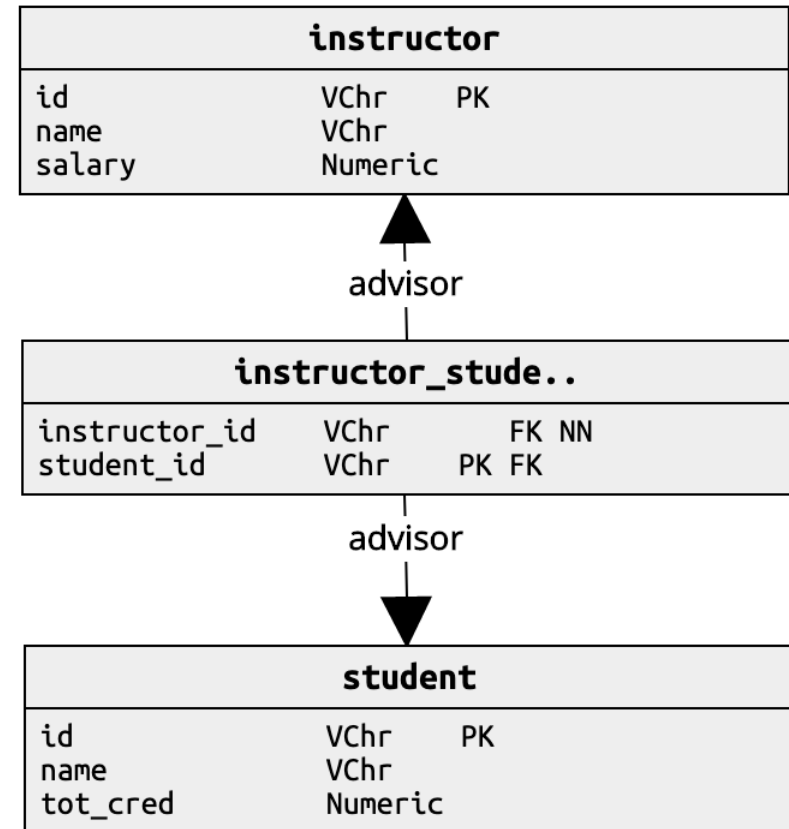
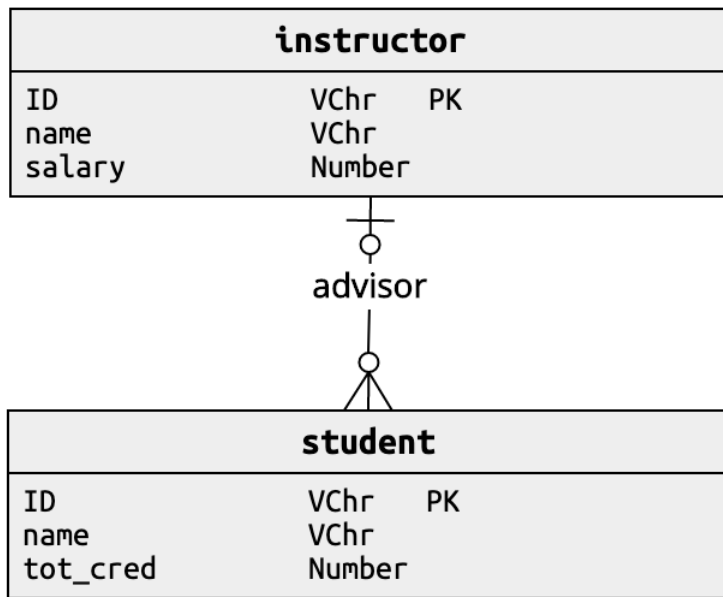


- Does this participation make any difference?

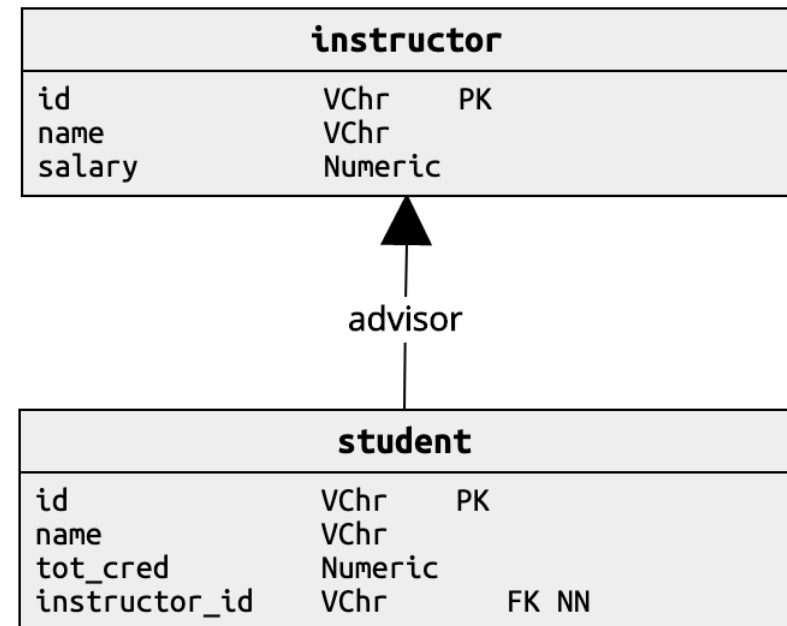
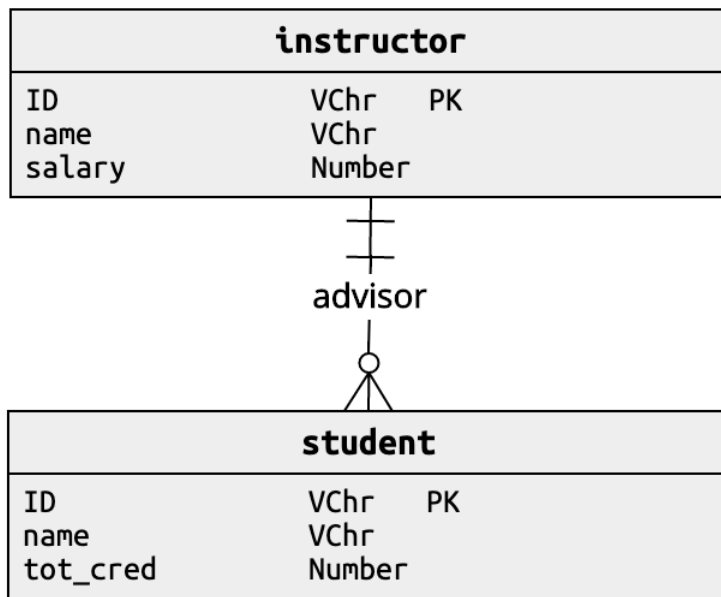


- So, why should it be represented in the E-R diagram?

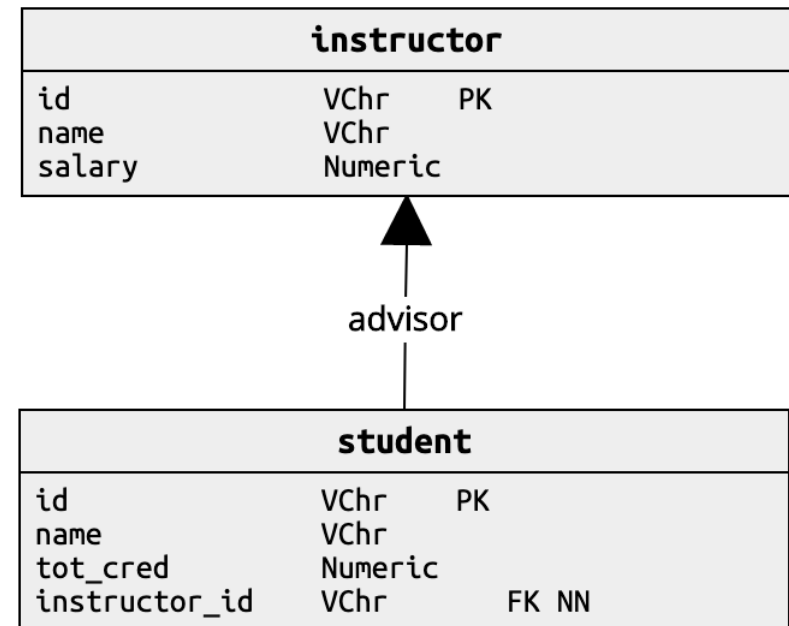
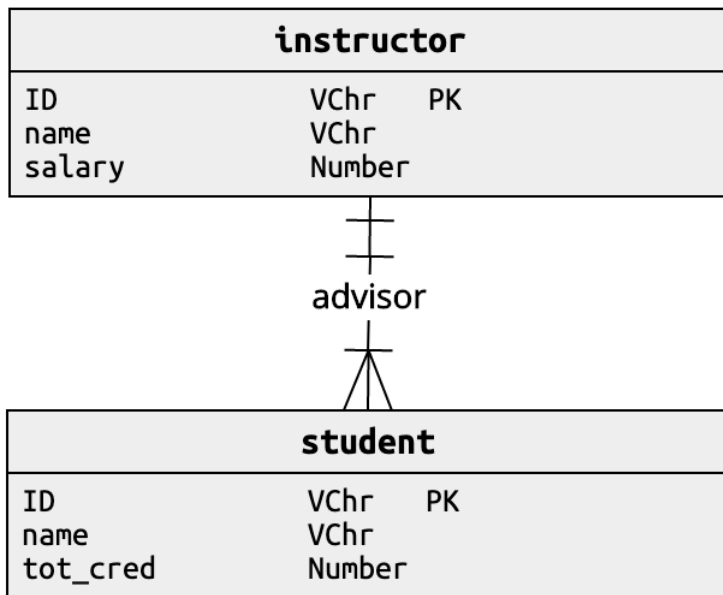
One-to-Many Relationship Sets: Examples



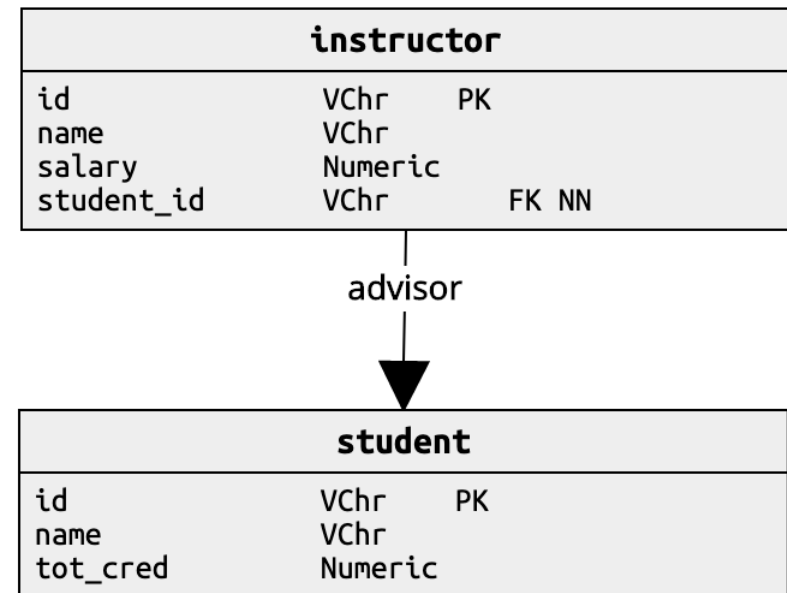
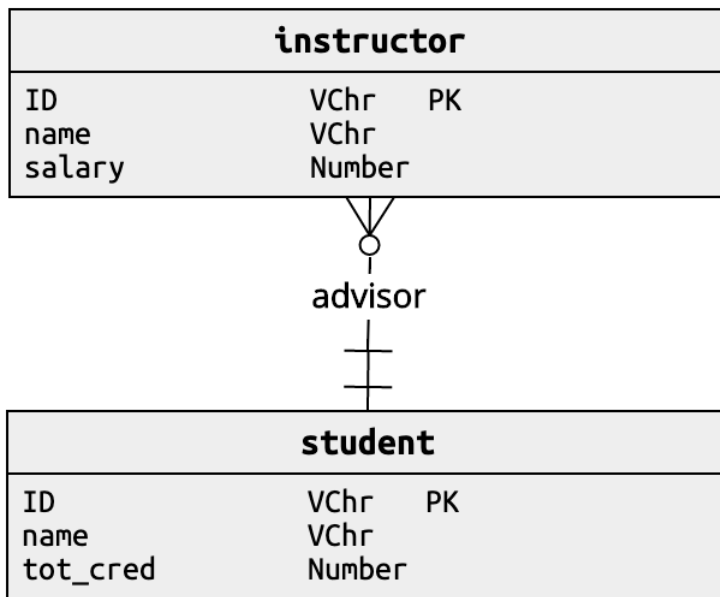
One-to-Many Relationship Sets: Examples



One-to-Many Relationship Sets: Examples




One-to-Many Relationship Sets: Examples


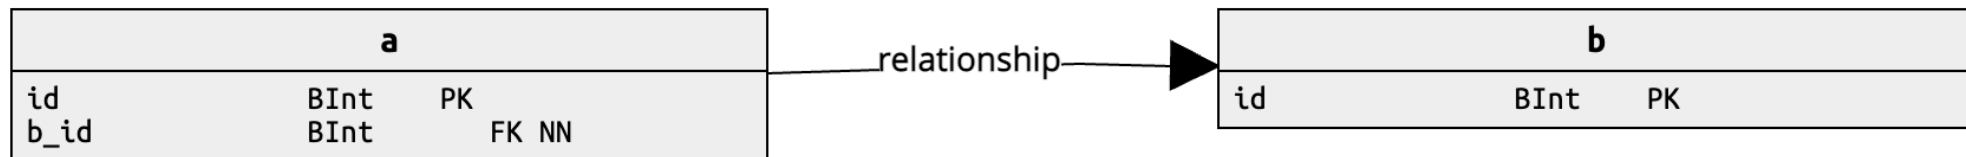


Many-to-Many Relationship Sets

- Let's assume a **many-to-many relationship set** with **partial participation on both sides**




- Can it be represented with two relations?

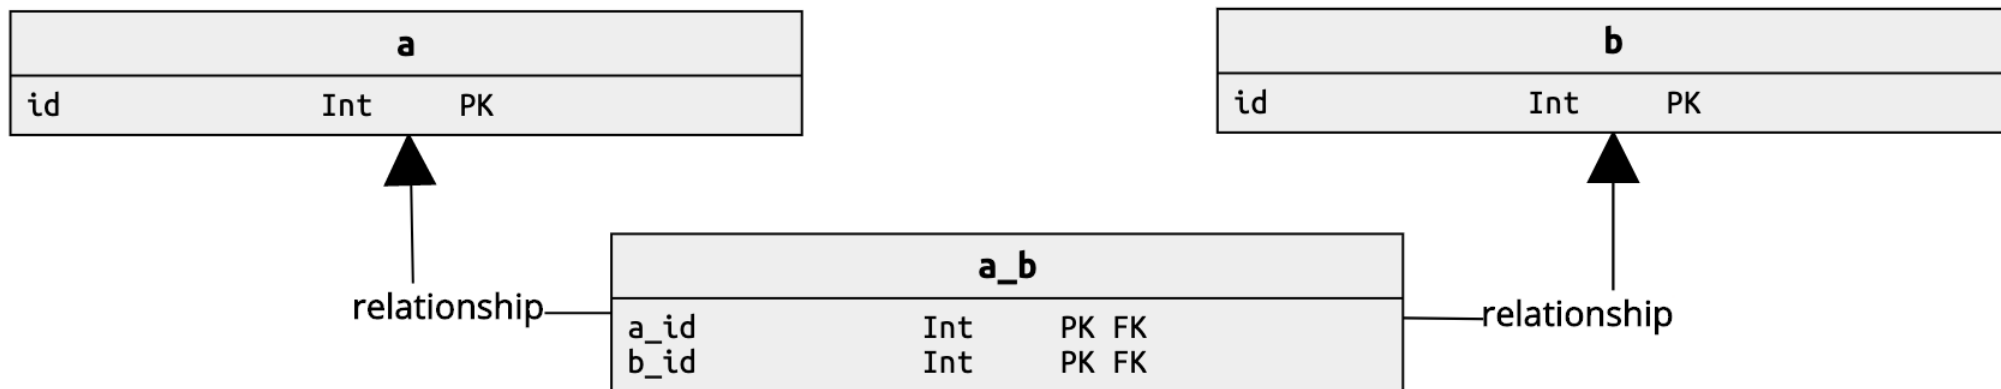


Many-to-Many Relationship Sets

- What about three relations?

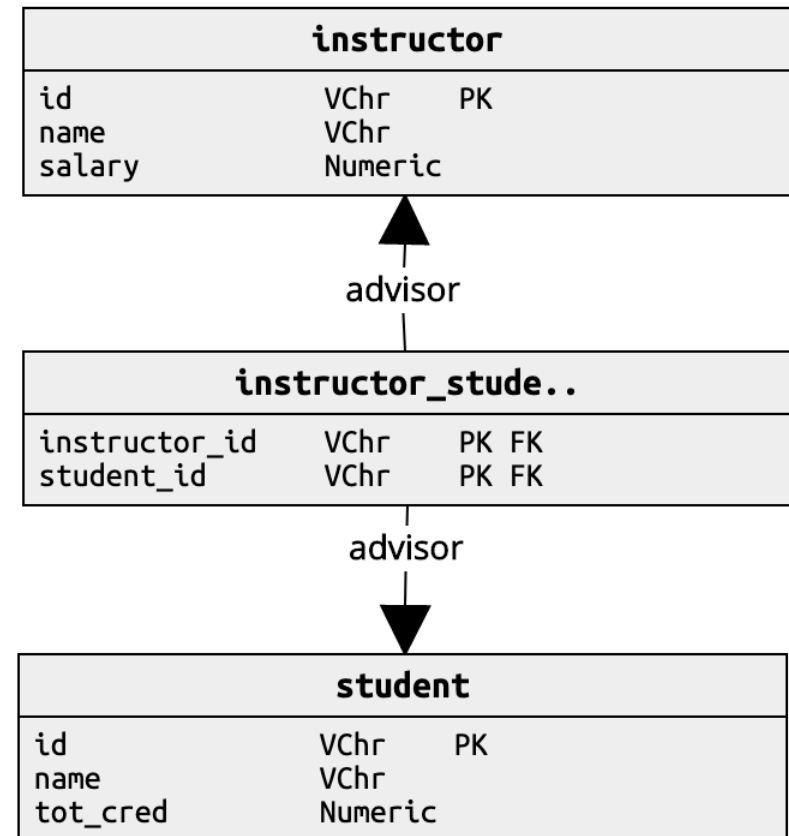
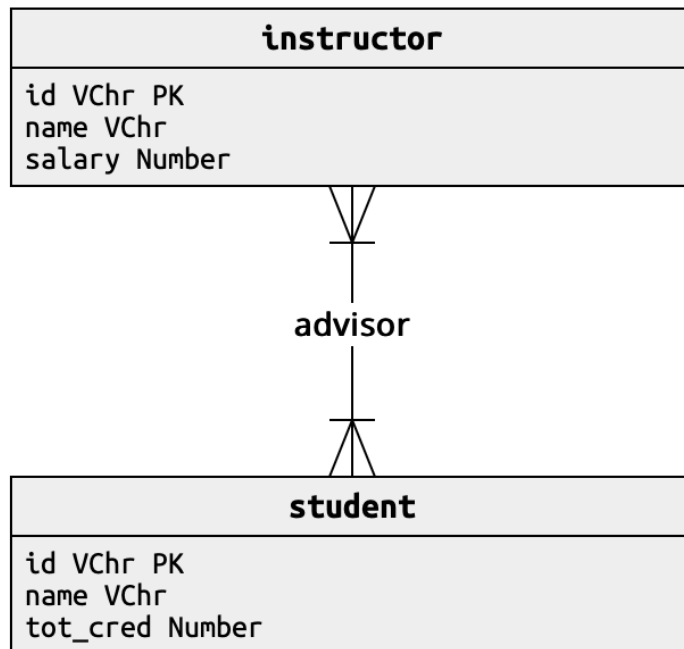


- What is the primary key of A_B ?
- Are there any foreign keys?



Many-to-Many Relationship Sets

- Does total participation make any difference in many-to-many relationships?
- So, why should total participation be represented in the E-R diagram?





Take-Away(s)

- Design Process: conceptual-design, logical-design, physical-design
- Entity-relationship (E-R) data model
- Entity and entity set
- Relationship and relationship set: binary, recursive
- Mapping cardinality: one-to-one, one-to-many, many-to-many
- Total and partial participation
- E-R diagram
- E-R diagram to relational schemas: different cases depending on the mapping cardinality and participation
- *onda* tool (and many others) can be used to support the process



Next Lesson(s)

- Weak Entity Sets
- Attributes of Relationship Sets
- *n*-ary Relationship Sets
- Extended E-R features
 - Specialization
 - Generalization
 - Attribute Inheritance
 - Completeness Constraints
- Typical design issues

Q&A



The seal of the University of Coimbra is partially visible on the left side of the slide. It is a circular emblem with a blue border containing the text 'SITATIS' at the bottom and 'CONIMBRIC' on the right. The central part of the seal depicts a figure, likely a saint or scholar, standing next to a large, ornate building with multiple towers and arches.

Databases

Database Design Using Entity-Relationship Model

João R. Campos

Bachelor in Informatics Engineering
Department of Informatics Engineering
University of Coimbra
2024/2025