

# Operating Systems 2024/2025

## T Class 05 – Deadlocks

Vasco Pereira (vasco@dei.uc.pt)

Dep. Eng. Informática da Faculdade de Ciências e Tecnologia da Universidade de Coimbra

### **operating system**

noun



the collection of software that directs a computer's operations, controlling and scheduling the execution of other programs, and managing storage, input/output, and communication resources.

Abbreviation: OS

Source: Dictionary.com

# Disclaimer

- This slides and notes are based on the companion material [Silberschatz13]. The original material can be found at:
  - <http://codex.cs.yale.edu/avi/os-book/OS9/slide-dir/>
- In some cases, material from [Stallings15] may also be used. The original material can be found at:
  - <http://williamstallings.com/OS/OS5e.html>
  - <http://williamstallings.com/OperatingSystems/>
- These slides also use materials from [McHoes2011] “Understanding Operating Systems, 6<sup>th</sup> Edition, Ann McIver McHoes and Ida M. Flynn”
- The respective copyrights belong to their owners.



**Note:** Some slides are also based on previous versions from Bruno Cabral, Paulo Marques and Luis Silva (Operating Systems classes of DEI-FCTUC).

# Deadlocks

## Outlines

- System model
- Characterization of a deadlock
- Deadlock prevention
- Deadlock avoidance
- Deadlock detection
- Recovery from a deadlock

# System model

- In multiprogramming different processes compete for a finite number of resources
  - (e.g.: CPU, I/O devices, mutexes, semaphores)
- How to use a resource?
  - 1 - **Request** - a process must request a resource before using it
    - If the resource is not available ,the process enters in a **waiting state**
  - 2 - **Use**
  - 3 - **Release** - resource must be released after being used
- To complete a task, a process may need different resources at the same time


# System model

- What is a deadlock?
  - Permanent blocking of a set of processes that either compete for system resources or communicate with each other.
- No efficient solution.
- Involves conflicting needs for resources by two or more processes.
- Typically, OSs do not provide deadlock-prevention facilities, and it remains the responsibility of programmers to ensure that they design deadlock-free programs.




# Deadlock characterization

- Why do deadlocks occur? A traffic example.



(a) Deadlock possible



(b) Deadlock

# Deadlock characterization

- Resources and deadlocks
  - 2 resources: A and B

Process P	Process Q
Get A	Get B
Get B	Get A
....	....
Release A	Release B
Release B	Release A

Execution trace with **deadlock**

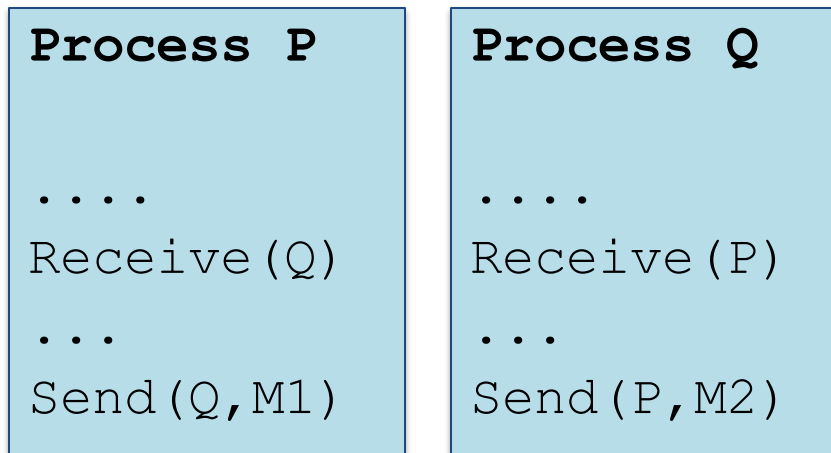
```
P - Get A
Q - Get B
P - Get B (block)
Q - Get A (block)
<DEADLOCK>
```

Execution trace without **deadlock**

```
P - Get A
P - Get B
Q - Get B (block)
P - ...
P - Release A
P - Release B
Q - Get B (unblock)
Q - Get A
Q - ...
Q - Release B
Q - Release A
```

# Deadlock characterization

- Operational semantics and deadlocks
  - If `receive()` is blocking...



Execution trace with **deadlock**

P - Receive(Q) <block>

Q - Receive(P) <block>


<DEADLOCK>



## Deadlock characterization



# Classical Problem: DINING PHILOSOPHERS

- Five philosophers sitting around a table
- Each has a plate of spaghetti
- There are forks between each pair of plates
- Philosophers need two forks to eat



# Deadlock characterization

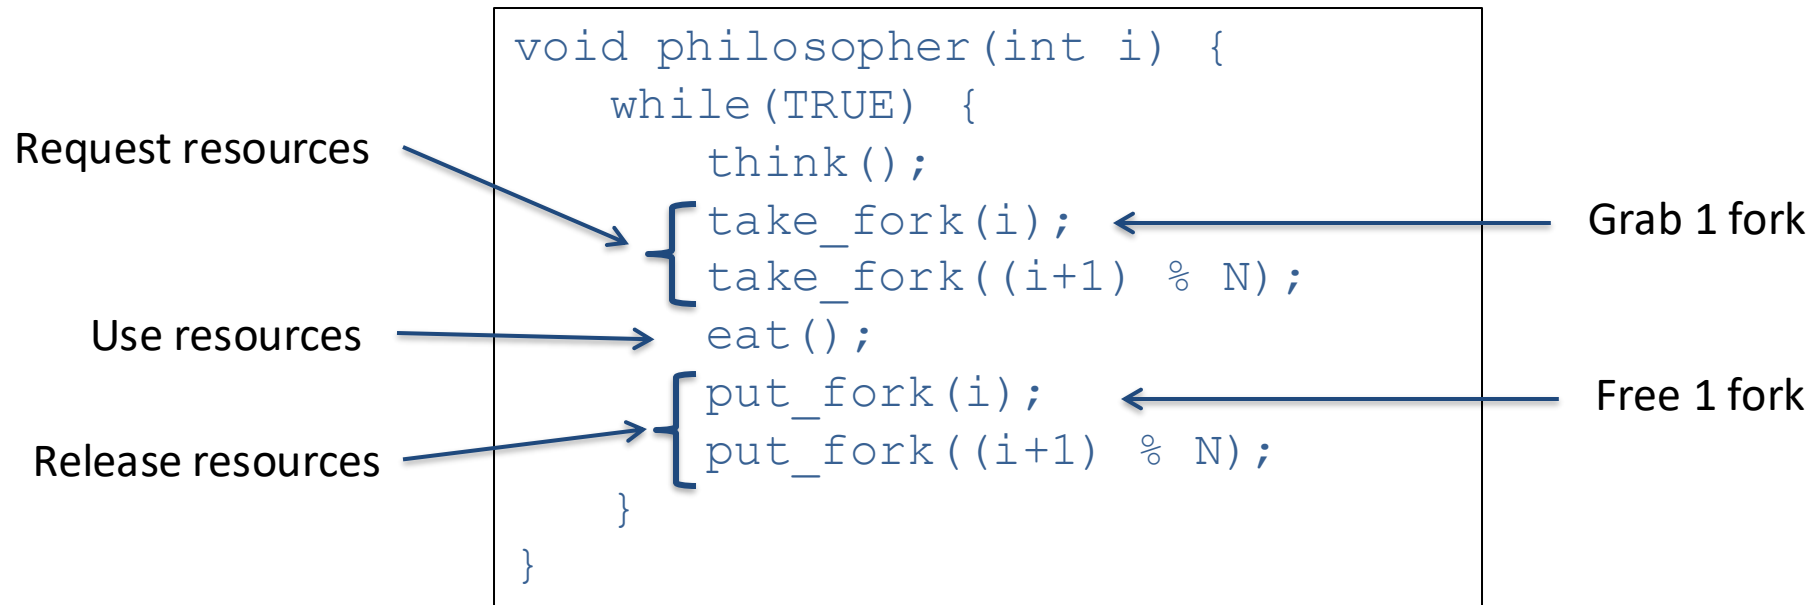
## DINING PHILOSOPHERS (2)



# Deadlock characterization

## DINING PHILOSOPHERS (3)


- A possible solution... is it good?



## Deadlock characterization

## DINING PHILOSOPHERS (4)

- Everyone grabs the first fork...




- **DEADLOCK...**

# Deadlock characterization

## Conditions for a Deadlock

- Deadlock arises if 4 conditions exist at the same time:
  - Mutual exclusion
    - Only one process may use a resource at a time
  - Hold-and-wait
    - A process holds a resource and is waiting for additional ones - a process requests all of its required resources at one time
  - No-preemption
    - Resources cannot be preempted, i.e., can only be released voluntarily.
  - Circular wait
    - A processes is waiting for a resource held by another process that in turn is waiting for a resource held by the first



# Deadlock characterization

## Resource-Allocation graph

### ■ Direct graph consisting of:


#### ■ Vertices

- Active processes (circles) - P
- Resources (rectangles) - R
  - Each dot represents an instance of the resource

#### ■ Directed edges

- $P_x \rightarrow R_y$ : Process  $P_x$  requested an instance of resource  $R_y$
- $R_y \rightarrow P_x$ : An instance of  $R_y$  has been allocated to  $P_x$


- It is used to describe the resource-allocation state at a specific time



## Deadlock characterization

## Resource-Allocation graph (1)


No  
deadlock



# Deadlock characterization

## Resource-Allocation graph (2)

Deadlock






## Deadlock characterization

## Resource-Allocation graph (3)

- With cycle but no deadlock



# Deadlock characterization

## Resource-Allocation graph (4)

- Resource-allocation graph analysis
  - If a graph contains no cycles  $\Rightarrow$  **no deadlock**.
  - If a graph contains a cycle ...
    - if only one instance per resource type, then **deadlock**.
    - if several instances per resource type, possibility of deadlock.

# Handling deadlocks

- Options for handling deadlocks in an OS:
  - 1 - **Prevent** deadlocks to ensure that the systems never enters a deadlock state
  - 2 - **Avoid** deadlocks to ensure that the systems never enters a deadlock state
  - 3 - **Allow** a deadlock state, **detect** it and **recover**
  - 4 - **Ignore** deadlocks
    - Is up to the programmers to handle deadlocks
    - Used in Linux and Windows

# Deadlock prevention

- Prevention is accomplished by ensuring that at least one of the 4 necessary conditions for a deadlock does not happen:
  - Mutual exclusion, hold-and-wait, no-preemption, circular wait
- Possible problems
  - Low device utilization
  - Reduced system throughput

# Deadlock prevention (2)

- Prevent Mutual exclusion

- Have sharable resources (e.g. read-only files)
  - However, not all resources can be sharable!

- Prevent Hold and Wait

- Guarantee that all the resources are requested and allocated before beginning execution
  - If all resources cannot be guaranteed release the ones allocated
- Or only request a resource after releasing the ones that are being hold
- Both options lead to low usage of resources and may lead to starvation of a process that uses different resources

# Deadlock prevention (3)

## ■ Prevent No-preemption

- If a process holding certain resources is denied a further request, that process must release its original resources.
- If a process requests a resource that is currently held by another process, the operating system may preempt the second process and require it to release its resources.

## ■ Prevent Circular Wait

- Prevented by defining a linear ordering of resource types

## Deadlock prevention

### Prevent a circular wait - Example

- Define a linear ordering of resources:

$R_1, R_2, R_3, \dots, R_n$

- If  $(i < j)$ :

Acquire ( $R_i$ )  
Acquire ( $R_j$ )

**OK**

Acquire ( $R_j$ )  
Acquire ( $R_i$ )


**NOT OK**

## Deadlock prevention

## Prevent a circular wait - Example

- Without linear ordering of resources

<b>Processo P1:</b> Sem_wait(A); Sem_wait(C); ..... Sem_signal(C); Sem_signal(A);	<b>Processo P2:</b> Sem_wait(B); Sem_wait(A); ..... Sem_signal(A); Sem_signal(B);	<b>Processo P3:</b> Sem_wait(C); Sem_wait(A); ..... Sem_signal(A); Sem_signal(C);	<b>Processo P4:</b> Sem_wait(B); Sem_wait(D); ..... Sem_signal(D); Sem_signal(B);
--	--	--	--



Supposing that P1, P2, P3 and P4 execute in sequence and that each executes one instruction...

**DEADLOCK**




# Deadlock prevention

## Prevent a circular wait - Example

- With linear ordering of resources

Processo P1:	Processo P2:	Processo P3:	Processo P4:
Sem_wait(A);	Sem_wait(A);	Sem_wait(A);	Sem_wait(B);
Sem_wait(C);	Sem_wait(B);	Sem_wait(C);	Sem_wait(D);
.....	.....	.....	.....
Sem_signal(C);	Sem_signal(B);	Sem_signal(C);	Sem_signal(D);
Sem_signal(A);	Sem_signal(A);	Sem_signal(A);	Sem_signal(B);




Supposing that P1, P2, P3 and P4 execute in sequence and that each executes one instruction...

After this point P1 and P4 start releasing resources...

**NO DEADLOCK**

# Back to the philosopher's example: can it be prevented?

- Proposal 1:
  - Instead of *take\_fork((i+1) % N)*, see if it is available first
  - If unavailable, **sleep** for a while and retry....
- Still no good....
  - what if everyone tries the left fork, waits, retries simultaneously, etc.
  - Everyone can run, but no one makes progress: **STARVATION**




# Back to the philosopher's example: can it be prevented?

## ■ Proposal 2:

- Before taking a fork, a philosopher locks a MUTEX
  - The philosopher can then take two forks, with no interference
  - When done eating, the forks are replaced and MUTEX is released
- 
- Nope.... only one philosopher can eat at a time...

```
void philosopher(int i) {  
    while(TRUE) {  
        think();  
        lock_mutex();  
        take_fork(i);  
        take_fork((i+1) % N);  
        eat();  
        put_fork(i);  
        put_fork((i+1) % N);  
        unlock_mutex();  
    }  
}
```


**C**



# Back to the philosopher's example: can it be prevented?

## ■ Solution:

```
void philosopher(int i) {  
    while(TRUE) {  
        think();  
        if(id %2 == 0) { // even number: left, right  
            take_fork(i);  
            take_fork((i+1) % N);  
        }  
        else{ // odd number: right, left  
            take_fork((i+1) % N);  
            take_fork(i);  
        }  
        eat();  
    }  
    ...  
}
```



**Forks are the resources! Take them by order!**

- Each even philosopher takes first the left fork
- Each odd philosopher takes first the right fork

# Deadlock avoidance

- How to avoid the occurrence of deadlocks?

- Require more information about the resources that are to be requested



- Approaches
    - Do not start a process if its demands might lead to a deadlock.
    - Do not grant an incremental resource request to a process if this allocation might lead to a deadlock.

# Deadlock avoidance

- Deadline avoidance algorithms dynamically examine the resource-allocation state to ensure that a circular-wait condition never happens
  - Is the state a **safe state**?

# Deadlock avoidance

## Safe state


- In a system with a fixed number of processes and resources:
  - The **state** of the system is the current allocation of resources to process.
  - A **safe sequence** exists when all processes may access their needed resources without resulting in a deadlock
  - A **safe state** is when there is at least one safe sequence of resource allocations that does not result in deadlock.
    - All processes can run to completion
    - The system can allocate resources to each process (up to its maximum) in some order and still avoid a deadlock.
  - Unsafe state is a state that is not safe (!).
    - An unsafe state may lead to a deadlock (**does not always result in a deadlock!**)

# Deadlock avoidance

## Resource-allocation-graph algorithm

### ■ Resource-allocation-graph algorithm

- Algorithm used when there is only one instance of each resource type
- The resource-allocation graph used has an additional edge – **claim edge**
  - A claim edge indicates that a process  $P_i$  may request a resource  $R_j$  in the future
  - It is depicted by a dashed line



- All resources must be claimed before the process starts executing
- When a process  $P_i$  requests resource  $R_j$ , the **claim edge**  $P_i \rightarrow R_j$  is converted to a **request edge**
- Requests are only granted if converting a **request edge** to an **assignment edge** does not result in a cycle in the graph



# Deadlock avoidance

## Banker's algorithm

### ■ Banker's Algorithm

- An algorithm that can be used in a banking system to allocate a limited amount of money to customers. The credit is given depending on the risk of the customer, so that the bank may continue to lend.
- Can be used with systems with multiple instances of each resource type
- A strategy of resource allocation denial
- Developed by Dijkstra

# Deadlock avoidance

## Banker's algorithm


- When a process enters the system, it declares the maximum number of instances of each resource type that it may need
- When resources are requested, the system must determine if their allocation leaves the system in a safe state
  - If it does -> resources are allocated
  - If it does not -> the process must wait until other process releases enough resources
- Support data structures
  - **Resource vector** – total amount of each resource in system
  - **Available vector** – amount of each resource not allocated
  - **Claim matrix** – total amount of resources that may be required by each process (stated before the process begins executing)
  - **Allocation matrix** – current allocation of resources to processes

# Deadlock avoidance

## Banker's algorithm

### ■ Determination of safe state

- Question: can any of the 4 processes run into completion with the available resources? I.e., is this a **safe state**?
  - What about P1?
  - What about P2?



(a) Initial state

# Deadlock avoidance

## Banker's algorithm

- Question: can any of the 4 processes run into completion with the available resources? (cont.)
  - P2 runs to completion!
  - What about P1?

P2 completed and released all resources. Now P1 can finish!

	R1	R2	R3
P1	3	2	2
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
6	2	3

Available vector V

(a) P2 runs to completion

# Deadlock avoidance

## Banker's algorithm

- Question: can any of the 4 processes run into completion with the available resources? (cont.)
  - P1 runs to completion!
  - What about P3?

P1 and P2 completed and released all resources. Now P3 can finish!

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
7	2	3

Available vector V

(a) P1 runs to completion

# Deadlock avoidance

## Banker's algorithm

- Question: can any of the 4 processes run into completion with the available resources?
  - P3 runs to completion!

The system was  
on a  
Safe State

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
9	3	4

Available vector V

P4 can also run  
until the end...

(a) P3 runs to completion

# Deadlock avoidance

## Banker's algorithm

### ■ **Deadlock Avoidance Strategy:**

(To assure that the system is always in a safe state)

- When a process makes a request for a set of resources, assume that the request is granted
- Update the system state accordingly and then determine if the result is a **SAFE STATE**.
  - If so, grant the request.
  - If not, block the process until it is safe to grant the request.

# Deadlock avoidance

## Banker's algorithm

### ■ Determination of an unsafe state

- Consider the following initial state:

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	1	0	2
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
1	1	2

Available vector V

(a) Initial state

- Assume that P1 makes a request: [1, 0, 1]
  - Should we give it those resources? I.e., does it lead to a safe state?



# Deadlock avoidance

## Banker's algorithm

- If request is granted, we will have:

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	2	0	1
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	1	2	1
P2	1	0	2
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
0	1	1

Available vector V

(b) P1 requests one unit each of R1 and R3

- Answer: No**, because in this state each process will need at least an additional unit of R1 and there are none available!
- Therefore, to **avoid** a **deadlock**, the request of P1 should be **denied**.
  - Maybe P1, P2, P3, P4 will not use any other R1 in the future, so the deadlock would not happen, but to avoid that possibility, the request is denied!

# Deadlock detection

- If a system does not use deadlock-prevention or avoidance algorithms a deadlock may occur!
- **Deadlock detection** strategies grant resources whenever possible and perform an algorithm to detect if a deadlock has occurred – does not prevent the deadlock, only detects it
  - Run periodically by the OS
- An algorithm to recover from the deadlock is also needed

## How to detect deadlocks?

# Deadlock detection Algorithm

## ■ Strategy

- Find processes whose resource requests can be satisfied
- Grant the resources and assume the process runs to completion and releases resources
  - This is an optimistic approach since the process may later request additional resources. However, what we want to detect is if a deadlock exists at this moment!
- Find another process and repeat
- The algorithm does not prevent deadlocks! It only determines if they exist at the time the algorithm runs.

# Deadlock detection Algorithm

- Uses 1 vector and 2 matrices
  - Available
    - Vector that indicates the number of available resources of each type
  - Allocation
    - Matrix that has the number of resources of each type currently allocated to each process
  - Request
    - Matrix that has the current requests of each process

# Deadlock detection Algorithm

1. Mark each process that has a row in the Allocation Matrix (**A**) of all zeros.
  - A process with no allocated resources does not participate in a deadlock.
2. Initialize a temporary vector **W** (Work) to be equal to the Available Vector (that lists resources from **1** to **m**).
3. Find an index **i** such that process **i** is currently unmarked and the **i<sup>th</sup>** row of the Request Matrix **Q** is less than or equal to **W**.
  - That is,  $Q_{ik} \leq W_k$  ( $k=1\dots m$ ). I.e. the process **i** can run with the available resources!
  - If no such row is found terminate the algorithm.
4. If such a row is found mark process **i** and add the corresponding row of the allocation matrix to **W**.
  - That is: set  $W_k = W_k + A_{ik}$  ( $k=1\dots m$ ). I.e. assume that on completion the process frees all the allocated resources.
  - Return to step 3.

# Deadlock detection Algorithm

- A **deadlock exists** if and only if there are unmarked processes at the end of the algorithm. The set of unmarked rows corresponds to the set of deadlocked processes.
  - These processes do not have sufficient resources available to fulfill their requests and cannot finish their execution

# Deadlock detection

## Algorithm - Example

	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request matrix Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation matrix A

R1	R2	R3	R4	R5
2	1	1	2	1

Resource vector

R1	R2	R3	R4	R5
0	0	0	0	1

Available vector

- **Mark**
- Set  $W = [0\ 0\ 0\ 0\ 1]$  (equals the Available vector)
- The request of process P3 is less than or equal to W, so **mark P3** and set:
  - $W = W + [0\ 0\ 0\ 1\ 0] = [0\ 0\ 0\ 1\ 1]$ 
    - This is the available vector after P3 finishes and releases all resources!
- No other unmarked process has a row in Q that is less than or equal to W. Therefore, terminate the algorithm.

P1 and P2 are unmarked, thus these processes are deadlocked!

# Recovery from a deadlock

- Abort all deadlocked processes.
- Successively abort deadlocked processes until deadlock no longer exists.
- Successively preempt resources until deadlock no longer exists.



# Deadlock Avoidance and Detection

## **EXERCISES**

# Exercise 1

- Consider a scenario where 4 processes are using 4 different system resources. The Request matrix (“Matriz de pedidos”), the Allocation matrix (“Matriz de recursos atribuídos”) and the Available resources vector (“Vector de recursos disponíveis”) are the following:

Allocation matrix A:

	R1	R2	R3	R4
P1	0	0	1	0
P2	0	0	1	1
P3	2	0	0	1
P4	0	1	2	0

Requests matrix Q:

	R1	R2	R3	R4
P1	1	0	2	0
P2	2	0	0	1
P3	1	0	1	1
P4	2	1	0	0

Available vector

R1	R2	R3	R4
2	1	0	0

Apply the deadlock detection algorithm and determine if there is a deadlock between the processes. Justify. (A binary answer is not enough, explain how the algorithm is applied).

**Solution:** There is a deadlock! Find out why!

# Exercise 2

- Consider the following scenario where 4 processes use 4 different system resources. Consider the given claim, allocated and maximum resources matrixes given (“Matriz de pedidos”, “Matriz de recursos alocados no momento” and “Matriz de recursos máximos”).
  - The maximum resource vector indicates the maximum number of resources of each type that exist in the system.
  - The maximum resources vector given is not the available resources vector!!

Requests matrix:

	R1	R2	R3	R4
P1	0	2	1	0
P2	0	2	0	0
P3	1	1	2	0
P4	0	0	2	1

Allocation resources matrix: Max resources vector:

	R1	R2	R3	R4
P1	1	1	0	1
P2	0	1	0	1
P3	0	0	2	1
P4	1	0	0	0

R1	R2	R3	R4
2	3	4	4

- Apply the deadlock detection algorithm and determine if there is a deadlock.
- If the max resources vector was [2 4 4 4] what would be the result?

# Exercise 3

- Consider a scenario with 5 processes (P1..P5) that use 4 different system resources (A,B,C,D) and the following allocation, claim and available matrices:

Claim

	A	B	C	D
P1	0	0	1	2
P2	1	7	5	0
P3	2	3	5	6
P4	0	6	5	2
P5	0	6	5	6

Allocation

	A	B	C	D
P1	0	0	1	2
P2	1	0	0	0
P3	1	3	5	4
P4	0	6	3	2
P5	0	0	1	4

Available

A	B	C	D
1	5	2	0

- Apply the Banker's algorithm and determine if the system is in a safe state.
- If P2 makes a new request (0,4,2,0), can the request be safely granted?

## Exercise 3 – solution (summary)

- Consider a scenario with 5 processes (P1..P5) that use 4 different system resources (A,B,C,D) and the following allocation, claim and available matrices:

Claim

	A	B	C	D
P1	0	0	1	2
P2	1	7	5	0
P3	2	3	5	6
P4	0	6	5	2
P5	0	6	5	6

Allocation

	A	B	C	D
P1	0	0	1	2
P2	1	0	0	0
P3	1	3	5	4
P4	0	6	3	2
P5	0	0	1	4

Available

A	B	C	D
1	5	2	0

- a) Apply the Banker's algorithm and determine if the system is in a safe state.

Is in a safe state

- b) If P2 makes a new request (0,4,2,0), can the request be safely granted?

Yes

# Exercise 4

- Consider a scenario with 5 processes (P1..P5) that use 4 shared different system resources (A,B,C,D) and the following allocation, claim (max resources needed by the process) and available matrices:

Allocation

	A	B	C	D
P1	0	0	1	2
P2	2	0	0	1
P3	0	1	1	2
P4	2	1	0	1
P5	0	0	2	1

Claim

	A	B	C	D
P1	1	1	1	2
P2	5	3	1	4
P3	3	1	2	3
P4	4	2	2	1
P5	1	1	3	3

Available

A	B	C	D
2	1	0	0

- Apply the Banker's algorithm and determine if the system is in a safe state or in an unsafe state.
- If P3 makes a new request (1,0,0,0), can the request be safely granted?


# Exercise 5

- Consider a scenario with 5 processes executing (P1..P5) and 3 different system resources (R1 to R3). The maximum number of resources provided by the system (is not the number of available ones now!) is: R1=10, R2=5, R3=7. Now, the system state is the following:

Resources in use				Maximum resources needed			
	R1	R2	R3		R1	R2	R3
P1	0	1	0		7	5	3
P2	2	0	0		3	2	2
P3	3	0	2		9	0	2
P4	2	1	1		2	2	2
P5	0	0	2		4	3	3

- Prove that the system is in a safe state.
- Process P5 makes the request (R1=3, R2=3, R3=0). Which will be the answer of the resource manager, if *Banker's Algorithm* is applied? Explain the process used to obtain the answer.

# References



- [Silberschatz13]
  - Chapter 7: Deadlocks



# Thank you! Questions?



*I keep six honest serving men. They taught me all I knew. Their names are What and Why and When and How and Where and Who.*  
—Rudyard Kipling